

Immersive_Data_Science_Project1_EDA

November 3, 2019

0.1 Final Project Submission

Please fill out: * Student name: Florian Fedrau

0.2 What to do

We are given the King County House Sales dataset and we want to predict the prices for houses.

0.3 Importing Python Libraries needed for the project

```
[1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
from pandas.plotting import scatter_matrix # data processing, scatter matrix
import statsmodels.api as sm # Python module that provides classes and
    ↪ functions for the estimation of many different
#statistical models, as well as for conducting statistical tests, and
    ↪ statistical data exploration
import matplotlib.pyplot as plt # plotting
import scipy.stats as stats
import sklearn.linear_model as linear_model
from sklearn.model_selection import train_test_split
import seaborn as sns
import statsmodels.api as sms
import statsmodels.formula.api as smf
from statsmodels.formula.api import ols
from mpl_toolkits.basemap import Basemap
sns.set()
%matplotlib inline
```

0.4 Import the data

```
[2]: df = pd.read_csv('King_County_House_prices_dataset.csv')
df.head()
```

```
[2]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
0	7129300520	10/13/2014	221900.0	3	1.00	1180	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	
2	5631500400	2/25/2015	180000.0	2	1.00	770	

3	2487200875	12/9/2014	604000.0	4	3.00	1960
4	1954400510	2/18/2015	510000.0	3	2.00	1680

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
0	5650	1.0	NaN	0.0	...	7	1180	
1	7242	2.0	0.0	0.0	...	7	2170	
2	10000	1.0	0.0	0.0	...	6	770	
3	5000	1.0	0.0	0.0	...	7	1050	
4	8080	1.0	0.0	0.0	...	8	1680	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
0	0.0	1955	0.0	98178	47.5112	-122.257	
1	400.0	1951	1991.0	98125	47.7210	-122.319	
2	0.0	1933	NaN	98028	47.7379	-122.233	
3	910.0	1965	0.0	98136	47.5208	-122.393	
4	0.0	1987	0.0	98074	47.6168	-122.045	

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

[5 rows x 21 columns]

0.5 Understanding the data

```
[3]: f = open("column_names.md", "r")
      print(f.read())
      f.close()
```

```
# Column Names and descriptions for Kings County Data Set
* **id** - unique identified for a house
* **dateDate** - house was sold
* **pricePrice** - is prediction target
* **bedroomsNumber** - of Bedrooms/House
* **bathroomsNumber** - of bathrooms/bedrooms
* **sqft_livingsquare** - footage of the home
* **sqft_lotsquare** - footage of the lot
* **floorsTotal** - floors (levels) in house
* **waterfront** - House which has a view to a waterfront
* **view** - Has been viewed
* **condition** - How good the condition is ( Overall )
* **grade** - overall grade given to the housing unit, based on King County
grading system
* **sqft_above** - square footage of house apart from basement
```

```

* **sqft_basement** - square footage of the basement
* **yr_built** - Built Year
* **yr_renovated** - Year when house was renovated
* **zipcode** - zip
* **lat** - Latitude coordinate
* **long** - Longitude coordinate
* **sqft_living15** - The square footage of interior housing living space for
the nearest 15 neighbors
* **sqft_lot15** - The square footage of the land lots of the nearest 15
neighbors

```

0.6 Cleaning the data

Lets get an overview about our data first.

```
[4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
id                21597 non-null int64
date              21597 non-null object
price             21597 non-null float64
bedrooms          21597 non-null int64
bathrooms          21597 non-null float64
sqft_living        21597 non-null int64
sqft_lot           21597 non-null int64
floors             21597 non-null float64
waterfront         19221 non-null float64
view               21534 non-null float64
condition          21597 non-null int64
grade              21597 non-null int64
sqft_above         21597 non-null int64
sqft_basement      21597 non-null object
yr_built           21597 non-null int64
yr_renovated       17755 non-null float64
zipcode            21597 non-null int64
lat                21597 non-null float64
long               21597 non-null float64
sqft_living15      21597 non-null int64
sqft_lot15         21597 non-null int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB

```

Sort for nan values

```
[5]: df.isna().sum().sort_values()
```

```
[5]: id          0
     long        0
     lat         0
     zipcode     0
     yr_built    0
     sqft_basement 0
     sqft_above  0
     grade       0
     sqft_living15 0
     condition   0
     floors      0
     sqft_lot    0
     sqft_living 0
     bathrooms   0
     bedrooms    0
     price       0
     date        0
     sqft_lot15  0
     view        63
     waterfront  2376
     yr_renovated 3842
     dtype: int64
```

We see that there are missing values in the columns view, waterfront and yr_renovated. Let's get an better understanding of the columns before we decide what to do with them.

First we take a look at the column waterfront

```
[6]: print("Unique values in waterfront:" + str(df.waterfront.unique()))
     print("Mean value of waterfront  : " + str(df.waterfront.mean()))
```

```
Unique values in waterfront:[nan  0.  1.]
Mean value of waterfront    :0.007595858696217679
```

Takeaway, basically it is an boolean saying if a house has waterfront. We can also see that most of the houses have no waterfront.

We asume that a waterfront is an unique selling point and every houseowner would give us the information if the house has one. So we decide to fill the "nan" with zeros.

```
[7]: df.waterfront.fillna(0, inplace =True)
     df.waterfront.unique()
```

```
[7]: array([0., 1.])
```

We continue with the next column view

```
[8]: print("Unique values in view : " + str(df.view.unique()))
     print("Mean value of view    : " + str(df.view.mean()))
     print("Number of Nan for view: " + str(df['view'].isna().sum()))
```

```
print("Unique count for view : ")
print(df.groupby('view')['id'].nunique())
```

```
Unique values in view : [ 0. nan  3.  4.  2.  1.]
Mean value of view      : 0.23386272870808952
Number of Nan for view: 63
Unique count for view :
view
0.0    19253
1.0     329
2.0     956
3.0     505
4.0     314
Name: id, dtype: int64
```

We can see that most of the houses have not been viewed. We have only 63 nan in the column and we can fill them with zeros.

```
[9]: df.view.fillna(0, inplace =True)
df.view.unique()
```

```
[9]: array([0., 3., 4., 2., 1.])
```

We continue with the column yr_renovated

```
[10]: print("Unique values in yr_renovated : " + str(df.yr_renovated.unique()))
print("Number of Nan for yr_renovated: " + str(df['yr_renovated'].isna().sum()))
print("Unique count for yr_renovated : ")
print(df.groupby('yr_renovated')['id'].nunique().head())
```

```
Unique values in yr_renovated : [  0. 1991.   nan 2002. 2010. 1992. 2013. 1994.
1978. 2005. 2003. 1984.
1954. 2014. 2011. 1983. 1945. 1990. 1988. 1977. 1981. 1995. 2000. 1999.
1998. 1970. 1989. 2004. 1986. 2007. 1987. 2006. 1985. 2001. 1980. 1971.
1979. 1997. 1950. 1969. 1948. 2009. 2015. 1974. 2008. 1968. 2012. 1963.
1951. 1962. 1953. 1993. 1996. 1955. 1982. 1956. 1940. 1976. 1946. 1975.
1964. 1973. 1957. 1959. 1960. 1967. 1965. 1934. 1972. 1944. 1958.]
Number of Nan for yr_renovated: 3842
Unique count for yr_renovated :
yr_renovated
0.0    16900
1934.0     1
1940.0     2
1944.0     1
1945.0     3
Name: id, dtype: int64
```

We assume that a renovation is a unique selling point and the houseowner would mention it. So we can fill the nan with a zero what means no renovation has taken place.

```
[11]: df.yr_renovated.fillna(0, inplace = True)
df['yr_renovated'].isna().sum()
```

```
[11]: 0
```

So lets take a look at our data:

```
[12]: df.isna().sum().sort_values()
```

```
[12]: id                0
long                0
lat                0
zipcode            0
yr_renovated        0
yr_built            0
sqft_basement        0
sqft_above          0
grade              0
sqft_living15        0
condition           0
waterfront          0
floors              0
sqft_lot            0
sqft_living          0
bathrooms           0
bedrooms            0
price               0
date                0
view                0
sqft_lot15          0
dtype: int64
```

We filled all of the missing data.

To be able to work with our data we want to have numeric values in our data. Lets figure out in what format our data is.

```
[13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
id                21597 non-null int64
date              21597 non-null object
price             21597 non-null float64
bedrooms          21597 non-null int64
bathrooms         21597 non-null float64
sqft_living        21597 non-null int64
sqft_lot          21597 non-null int64
```

```

floors          21597 non-null float64
waterfront      21597 non-null float64
view            21597 non-null float64
condition        21597 non-null int64
grade           21597 non-null int64
sqft_above      21597 non-null int64
sqft_basement   21597 non-null object
yr_built        21597 non-null int64
yr_renovated    21597 non-null float64
zipcode         21597 non-null int64
lat             21597 non-null float64
long            21597 non-null float64
sqft_living15   21597 non-null int64
sqft_lot15      21597 non-null int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB

```

So the columns date and sqft_basement are objects. We want to change that. First we take a look at sqft_basement. I would expect an numeric value. Lets explore sqft_basement

```
[14]: print(df.groupby('sqft_basement')['id'].nunique())
```

```

sqft_basement
0.0          12718
10.0           1
100.0          42
1000.0         146
1008.0          1
1010.0          62
1020.0          51
1024.0           1
1030.0          44
1040.0          54
1050.0          41
1060.0          58
1070.0          51
1080.0          31
1090.0          32
110.0          18
1100.0          78
1110.0          35
1120.0          43
1130.0          30
1135.0           1
1140.0          28
1150.0          26
1160.0          26
1170.0          30

```

```

1180.0    28
1190.0    24
120.0     53
1200.0    68
1210.0    18
...
80.0      20
800.0    201
810.0     55
820.0     62
830.0     56
840.0     82
850.0     69
860.0     79
861.0      1
862.0      1
870.0     48
875.0      1
880.0     69
890.0     52
90.0      21
900.0    141
906.0      1
910.0     69
915.0      1
920.0     65
930.0     41
935.0      1
940.0     71
946.0      1
950.0     62
960.0     65
970.0     44
980.0     55
990.0     51
?         454

```

```
Name: id, Length: 304, dtype: int64
```

We see the unique `sqft_basement` values and detect that in 454 rows is a “?”. Thats why the datatype of the column is an object. We definetly want to change that.

So lets try to find some information about the houses with basement, without basement and we don't know about the basement.

```
[15]: df[['sqft_basement', 'sqft_above', 'sqft_living']].head()
```

```
[15]:   sqft_basement  sqft_above  sqft_living
0           0.0         1180         1180
1          400.0         2170         2570
```


2	0.0	770	770
3	910.0	1050	1960
4	0.0	1680	1680

```
[16]: basement = df[['sqft_basement', 'sqft_above', 'sqft_living']]
      basement['sqft_living-above'] = basement['sqft_living'] - basement['sqft_above']
      basement
```

/Users/flori/anaconda3/envs/nf/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
[16]:
```

	sqft_basement	sqft_above	sqft_living	sqft_living-above
0	0.0	1180	1180	0
1	400.0	2170	2570	400
2	0.0	770	770	0
3	910.0	1050	1960	910
4	0.0	1680	1680	0
5	1530.0	3890	5420	1530
6	?	1715	1715	0
7	0.0	1060	1060	0
8	730.0	1050	1780	730
9	0.0	1890	1890	0
10	1700.0	1860	3560	1700
11	300.0	860	1160	300
12	0.0	1430	1430	0
13	0.0	1370	1370	0
14	0.0	1810	1810	0
15	970.0	1980	2950	970
16	0.0	1890	1890	0
17	0.0	1600	1600	0
18	?	1200	1200	0
19	0.0	1250	1250	0
20	760.0	860	1620	760
21	720.0	2330	3050	720
22	0.0	2270	2270	0
23	0.0	1070	1070	0
24	0.0	2450	2450	0
25	0.0	1710	1710	0
26	700.0	1750	2450	700
27	0.0	1400	1400	0
28	730.0	790	1520	730

29	0.0	2570	2570	0
...
21567	0.0	710	710	0
21568	320.0	940	1260	320
21569	0.0	1870	1870	0
21570	0.0	1430	1430	0
21571	0.0	1520	1520	0
21572	190.0	1020	1210	190
21573	0.0	2540	2540	0
21574	1800.0	3110	4910	1800
21575	0.0	2770	2770	0
21576	0.0	1190	1190	0
21577	0.0	4170	4170	0
21578	0.0	2500	2500	0
21579	50.0	1480	1530	50
21580	0.0	3600	3600	0
21581	?	3410	3410	0
21582	0.0	3118	3118	0
21583	0.0	3990	3990	0
21584	0.0	4470	4470	0
21585	0.0	1425	1425	0
21586	0.0	1500	1500	0
21587	0.0	2270	2270	0
21588	0.0	1490	1490	0
21589	0.0	2520	2520	0
21590	910.0	2600	3510	910
21591	130.0	1180	1310	130
21592	0.0	1530	1530	0
21593	0.0	2310	2310	0
21594	0.0	1020	1020	0
21595	0.0	1600	1600	0
21596	0.0	1020	1020	0

[21597 rows x 4 columns]

Seems like the sqft_basement is just calculated by subtracting sqft_above from sqft_living

```
[17]: df['sqft_basement'] = basement['sqft_living'] - basement['sqft_above']
```

So sqft_living is just the sum of sqft_aboce and sqft_basement we can just drop it.

```
[18]: df.drop('sqft_living', axis = 1, inplace = True)
```

```
[19]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 20 columns):
```

```

id                21597 non-null int64
date              21597 non-null object
price             21597 non-null float64
bedrooms         21597 non-null int64
bathrooms        21597 non-null float64
sqft_lot         21597 non-null int64
floors           21597 non-null float64
waterfront       21597 non-null float64
view             21597 non-null float64
condition        21597 non-null int64
grade            21597 non-null int64
sqft_above       21597 non-null int64
sqft_basement    21597 non-null int64
yr_built         21597 non-null int64
yr_renovated     21597 non-null float64
zipcode          21597 non-null int64
lat              21597 non-null float64
long             21597 non-null float64
sqft_living15    21597 non-null int64
sqft_lot15       21597 non-null int64
dtypes: float64(8), int64(11), object(1)
memory usage: 3.3+ MB

```

Now we have to look at the date column. To be able to work with it we have to change it to a numerical value. Therefore we want to separate between spring, summer, autumn and winter. We assume that the winter starts at december 1st, spring starts at march 1st, summer starts at june 1st and autumn at september 1st.

```
[20]: df['date'] = pd.to_datetime(df['date'])
```

```
[21]: dates = []
for i in range(len(df)):
    if(df['date'][i].month == 12 or df['date'][i].month == 1 or df['date'][i].
    month == 2):
        dates.append('winter')
    elif(df['date'][i].month == 3 or df['date'][i].month == 4 or df['date'][i].
    month == 5):
        dates.append('spring')
    elif(df['date'][i].month == 6 or df['date'][i].month == 7 or df['date'][i].
    month == 8):
        dates.append('summer')
    else:
        dates.append('autumn')
```

```
[22]: dates_dummies = pd.get_dummies(dates, drop_first=True)
dates_dummies.head()
```

```
[22]:      spring  summer  winter
0         0         0         0
1         0         0         1
2         0         0         1
3         0         0         1
4         0         0         1
```

We generated our dummy variables. Now we want to join our dummy variables with our dataset

```
[23]: df2 = pd.concat([df, dates_dummies], axis = 1)
df2.head()
```

```
[23]:      id      date      price  bedrooms  bathrooms  sqft_lot  floors  \
0  7129300520  2014-10-13  221900.0         3         1.00     5650     1.0
1  6414100192  2014-12-09  538000.0         3         2.25     7242     2.0
2  5631500400  2015-02-25  180000.0         2         1.00    10000     1.0
3  2487200875  2014-12-09  604000.0         4         3.00     5000     1.0
4  1954400510  2015-02-18  510000.0         3         2.00     8080     1.0

      waterfront  view  condition  ...  yr_built  yr_renovated  zipcode  \
0           0.0   0.0          3  ...    1955           0.0    98178
1           0.0   0.0          3  ...    1951          1991.0    98125
2           0.0   0.0          3  ...    1933           0.0    98028
3           0.0   0.0          5  ...    1965           0.0    98136
4           0.0   0.0          3  ...    1987           0.0    98074

      lat      long  sqft_living15  sqft_lot15  spring  summer  winter
0  47.5112 -122.257         1340         5650         0         0         0
1  47.7210 -122.319         1690         7639         0         0         1
2  47.7379 -122.233         2720         8062         0         0         1
3  47.5208 -122.393         1360         5000         0         0         1
4  47.6168 -122.045         1800         7503         0         0         1
```

[5 rows x 23 columns]

Now we want to drop the month and day part of the date, so we just have the year left.

```
[24]: df2['date'] = pd.DatetimeIndex(df2['date']).year
```

```
[25]: df2.head()
```

```
[25]:      id  date      price  bedrooms  bathrooms  sqft_lot  floors  \
0  7129300520  2014  221900.0         3         1.00     5650     1.0
1  6414100192  2014  538000.0         3         2.25     7242     2.0
2  5631500400  2015  180000.0         2         1.00    10000     1.0
3  2487200875  2014  604000.0         4         3.00     5000     1.0
4  1954400510  2015  510000.0         3         2.00     8080     1.0
```

	waterfront	view	condition	...	yr_built	yr_renovated	zipcode	\
0	0.0	0.0	3	...	1955	0.0	98178	
1	0.0	0.0	3	...	1951	1991.0	98125	
2	0.0	0.0	3	...	1933	0.0	98028	
3	0.0	0.0	5	...	1965	0.0	98136	
4	0.0	0.0	3	...	1987	0.0	98074	


	lat	long	sqft_living15	sqft_lot15	spring	summer	winter
0	47.5112	-122.257	1340	5650	0	0	0
1	47.7210	-122.319	1690	7639	0	0	1
2	47.7379	-122.233	2720	8062	0	0	1
3	47.5208	-122.393	1360	5000	0	0	1
4	47.6168	-122.045	1800	7503	0	0	1

[5 rows x 23 columns]

Create new column "last_building_phase" include the max of yr_built and yr_renovated

```
[26]: df2["last_building_phase"] = df2[["yr_built", "yr_renovated"]].max(axis=1)
```


```
[27]: sns.scatterplot(x="last_building_phase", y="price", data=df2);
```



There is a small relationship between price and last_building_phase.

Lets plot some data

```
[28]: corr = df2.corr()
f, ax = plt.subplots(figsize = (18,18))
sns.heatmap(data=corr, center = 0, cmap="coolwarm", annot=True);
```



There are variables that have a correlation with price:

- bedrooms 0.31
- bathrooms 0.53
- sqft-lot 0.09
- floors 0.26
- waterfront 0.26
- view 0.39
- grade 0.67

```

-sqft_above 0.31
-sqft_basement 0.32
-year_renovated 0.12
-lat 0.31
-sqft_living 0.59
-last_building_phase 0.098

```

Lets plot the location and the price:

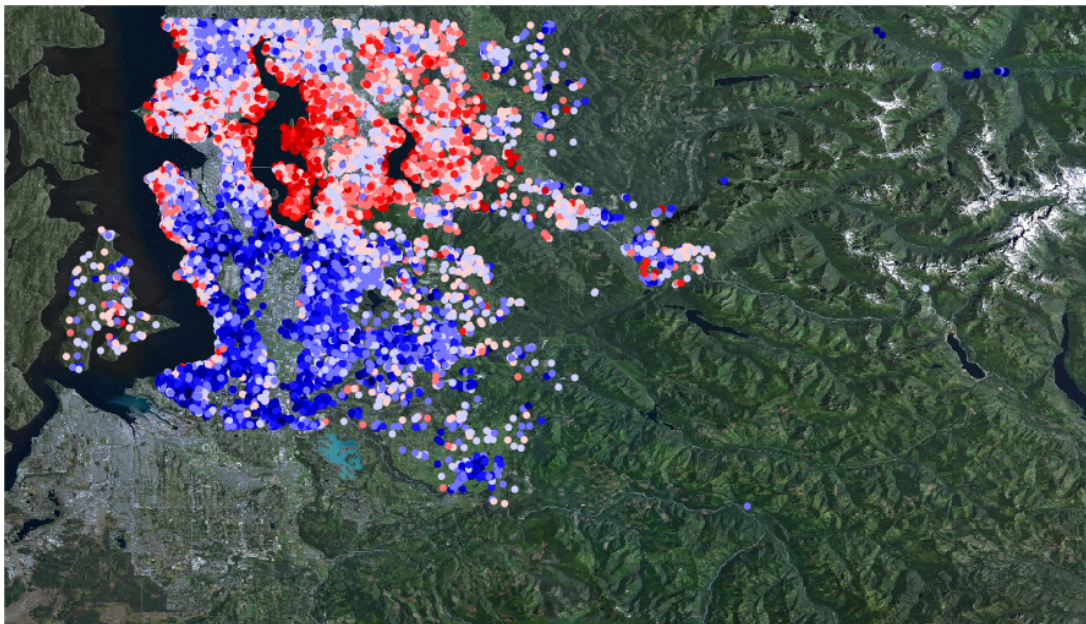
```

[29]: # Extract the data we're interested in
lat = df['lat'].values
lon = df['long'].values
bins = [0, 150000, 200000, 250000, 300000, 400000, 550000, 750000, 1000000,
↪1500000, 2500000, 5000000, 10000000]
c = pd.cut(df['price'], bins, labels = range(12))
# 1. Draw the map background
fig = plt.figure(figsize=(16, 16))
m = Basemap(projection='merc', resolution='h', llcrnrlon = -122.6, llcrnrlat =
↪47, urcrnrlon = -121.2, urcrnrlat = 47.8,
            lon_0= -121.7,
            lat_0= 47.4,
            epsg=4269)
m.arcgisimage(service='ESRI_Imagery_World_2D', xpixels = 2000)

# 2. scatter city data, with color reflecting population
# and size reflecting area
lons, lats = m(lon, lat)
m.scatter(lons, lats, zorder=1, linewidths=0.07, cmap = 'seismic', c = c)

plt.show()

```



We see that there is a correlation between location and price.

Let's take the squares of the plot and group our houses into them

```
[30]: lat = []
for i in (range(len(df))):
    if df['lat'][i] < 47.2:
        lat.append('lat_47.2')
    elif df['lat'][i] < 47.3:
        lat.append('lat_47.3')
    elif df['lat'][i] < 47.4:
        lat.append('lat_47.4')
    elif df['lat'][i] < 47.5:
        lat.append('lat_47.5')
    elif df['lat'][i] < 47.6:
        lat.append('lat_47.6')
    elif df['lat'][i] < 47.7:
        lat.append('lat_47.7')
    else:
        lat.append('lat_47.8')
```

```
[31]: long = []
for i in (range(len(df))):
    if df['long'][i] < -122.4:
        long.append('long_-122.4')
    elif df['long'][i] < -122.2:
        long.append('long_-122.2')
```



```

elif df['long'][i] < -122:
    long.append('long_-122')
elif df['long'][i] < -121.8:
    long.append('long_-121.8')
elif df['long'][i] < -121.6:
    long.append('long_-121.6')
elif df['long'][i] < -121.4:
    long.append('long_-121.4')
else:
    long.append('long_-121.2')

```

```

[32]: lat_dummies = pd.get_dummies(lat, drop_first=True)
      long_dummies = pd.get_dummies(long, drop_first=True)
      long_dummies.head()

```

```

[32]:   long_-121.4  long_-121.6  long_-121.8  long_-122  long_-122.2  long_-122.4
0           0           0           0           0           1           0
1           0           0           0           0           1           0
2           0           0           0           0           1           0
3           0           0           0           0           1           0
4           0           0           0           1           0           0

```

```

[33]: lat_long = pd.concat([lat_dummies, long_dummies], axis = 1)
      lat_long.head()

```

```

[33]:   lat_47.3  lat_47.4  lat_47.5  lat_47.6  lat_47.7  lat_47.8  long_-121.4  \
0           0           0           0           1           0           0           0
1           0           0           0           0           0           1           0
2           0           0           0           0           0           1           0
3           0           0           0           1           0           0           0
4           0           0           0           0           1           0           0

      long_-121.6  long_-121.8  long_-122  long_-122.2  long_-122.4
0           0           0           0           1           0
1           0           0           0           1           0
2           0           0           0           1           0
3           0           0           0           1           0
4           0           0           1           0           0

```

```

[34]: df3 = pd.concat([df2, lat_long], axis = 1)

```

0.7 Model a linear Regression

```

[35]: X = df3[['bedrooms', 'bathrooms',
              'sqft_lot', 'floors', 'waterfront', 'view', 'grade',
              'sqft_above', 'sqft_basement', 'lat_47.8', 'lat_47.3', 'lat_47.4',
              'lat_47.5', 'lat_47.6', 'lat_47.7', 'long_-122.4', 'long_-121.4',

```

```

        'long_-121.6', 'long_-121.8', 'long_-122', 'long_-122.2']]
y = df3[['price']]
X = sms.add_constant(X)
model = sms.OLS(y, X)
results = model.fit()
results.summary()

```

```

[35]: <class 'statsmodels.iolib.summary.Summary'>
      """

```

```

                                OLS Regression Results
=====
Dep. Variable:                price    R-squared:                0.727
Model:                        OLS      Adj. R-squared:            0.727
Method:                    Least Squares    F-statistic:                2740.
Date:                Sun, 03 Nov 2019    Prob (F-statistic):            0.00
Time:                10:10:31    Log-Likelihood:            -2.9336e+05
No. Observations:                21597    AIC:                5.868e+05
Df Residuals:                21575    BIC:                5.869e+05
Df Model:                21
Covariance Type:                nonrobust
=====
=
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
-
const                -4.409e+05    6.26e+04     -7.039     0.000    -5.64e+05
-3.18e+05
bedrooms            -2.451e+04    1806.001    -13.570     0.000    -2.8e+04
-2.1e+04
bathrooms             1.454e+04    2936.901      4.952     0.000     8786.795
2.03e+04
sqft_lot              0.1172         0.034       3.485     0.000         0.051
0.183
floors              -7.445e+04    3338.798    -22.299     0.000    -8.1e+04
-6.79e+04
waterfront           6.575e+05    1.73e+04     38.000     0.000     6.24e+05
6.91e+05
view                 5.832e+04    1999.164     29.171     0.000     5.44e+04
6.22e+04
grade                6.656e+04    1941.184     34.291     0.000     6.28e+04
7.04e+04
sqft_above           230.2617         3.150      73.103     0.000       224.088
236.436
sqft_basement        155.4998         4.087      38.044     0.000       147.488
163.511
lat_47.8             1.164e+05    2.17e+04      5.370     0.000     7.39e+04

```

```

1.59e+05
lat_47.3      -5.139e+04  2.25e+04   -2.281      0.023   -9.56e+04
-7235.817
lat_47.4      -4.009e+04  2.17e+04   -1.845      0.065   -8.27e+04
2489.904
lat_47.5       1.744e+04  2.17e+04    0.802      0.422   -2.52e+04
6.01e+04
lat_47.6       1.766e+05  2.16e+04    8.192      0.000   1.34e+05
2.19e+05
lat_47.7       3.217e+05  2.16e+04   14.869      0.000   2.79e+05
3.64e+05
long_-122.4    6.164e+04  5.9e+04     1.046      0.296   -5.39e+04
1.77e+05
long_-121.4   -1.014e+05  9.74e+04   -1.041      0.298   -2.92e+05
8.96e+04
long_-121.6    2.389e+04  5.95e+04    0.401      0.688   -9.28e+04
1.41e+05
long_-121.8   -9.844e+04  5.83e+04   -1.689      0.091   -2.13e+05
1.58e+04
long_-122     -2.936e+04  5.81e+04   -0.506      0.613   -1.43e+05
8.45e+04
long_-122.2    4.69e+04  5.8e+04     0.809      0.419   -6.68e+04
1.61e+05
=====
Omnibus:                18904.678   Durbin-Watson:                1.990
Prob(Omnibus):           0.000   Jarque-Bera (JB):            2171070.288
Skew:                    3.701   Prob(JB):                     0.00
Kurtosis:                51.558   Cond. No.                     5.31e+06
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.31e+06. This might indicate that there are strong multicollinearity or other numerical problems.

"""

R-squared is 0.727 not bad at all.

0.8 Improving our approach

We want to introduce dummy variables for the zipcode, because the zipcode is an integer but is actual a categorial variable.

First we plot the zipcode and try to get some insights

```

[36]: # Extract the data we're interested in
lat = df['lat'].values

```

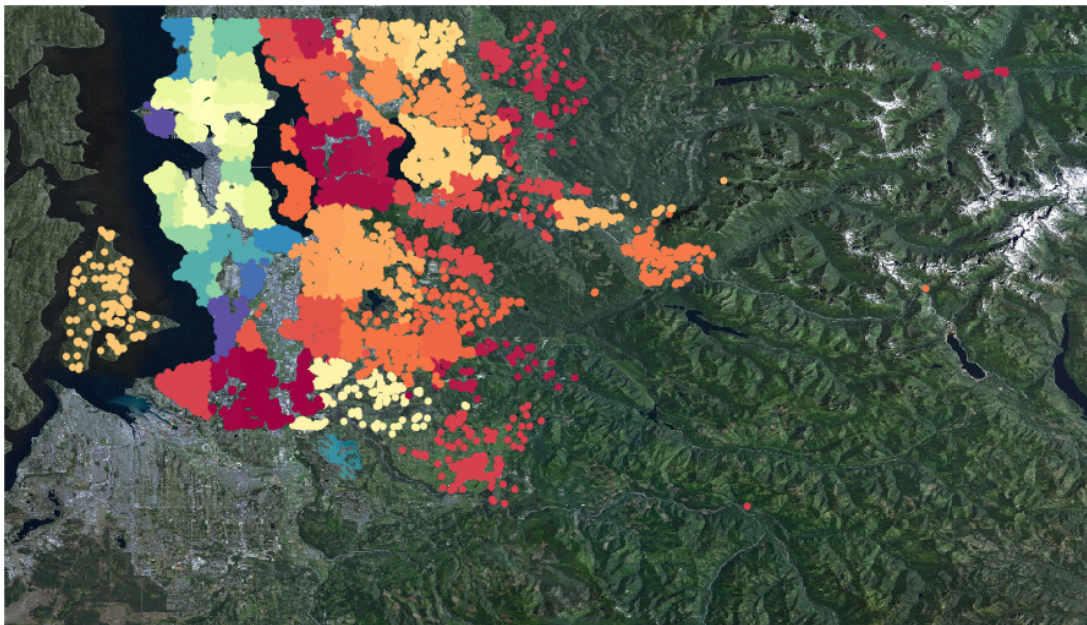
```

lon = df['long'].values
# 1. Draw the map background
fig = plt.figure(figsize=(16, 16))
m = Basemap(projection='merc', resolution='h', llcrnrlon = -122.6, llcrnrlat = 47.4,
            urcrnrlon = -121.2, urcrnrlat = 47.8,
            lon_0= -121.7,
            lat_0= 47.4,
            epsg=4269)
m.arcgisimage(service='ESRI_Imagery_World_2D', xpixels = 2000)

# 2. scatter city data, with color reflecting population
# and size reflecting area
lons, lats = m(lon, lat)
m.scatter(lons, lats, zorder=1, linewidths=0.1, c = df['zipcode'].values, cmap=
        'Spectral')

plt.show()

```



Seems to work quite good and should be more exact than just the dummies for longitude and latitude.

```

[37]: model = ols('price ~ C(zipcode) + bedrooms + bathrooms + sqft_lot + floors +
        waterfront + view + grade + sqft_above + sqft_basement +
        last_building_phase', data=df3).fit()

model.summary()

```

```
[37]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
                                OLS Regression Results
=====
Dep. Variable:                price    R-squared:                0.804
Model:                        OLS      Adj. R-squared:           0.804
Method:                        Least Squares    F-statistic:            1119.
Date:                          Sun, 03 Nov 2019    Prob (F-statistic):      0.00
Time:                          10:10:54    Log-Likelihood:         -2.8978e+05
No. Observations:              21597    AIC:                    5.797e+05
Df Residuals:                  21517    BIC:                    5.804e+05
Df Model:                       79
Covariance Type:                nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept                1.056e+06    1.13e+05     9.346     0.000     8.34e+05
1.28e+06
C(zipcode) [T.98002]    3.497e+04    1.44e+04     2.429     0.015     6747.289
6.32e+04
C(zipcode) [T.98003]   -1.719e+04    1.3e+04     -1.324     0.185    -4.26e+04
8253.120
C(zipcode) [T.98004]    7.868e+05    1.27e+04    62.025     0.000     7.62e+05
8.12e+05
C(zipcode) [T.98005]    3.104e+05    1.53e+04    20.244     0.000     2.8e+05
3.41e+05
C(zipcode) [T.98006]    2.772e+05    1.14e+04    24.239     0.000     2.55e+05
3e+05
C(zipcode) [T.98007]    2.513e+05    1.62e+04    15.488     0.000     2.19e+05
2.83e+05
C(zipcode) [T.98008]    2.538e+05    1.3e+04    19.529     0.000     2.28e+05
2.79e+05
C(zipcode) [T.98010]     7.7e+04    1.84e+04     4.174     0.000     4.08e+04
1.13e+05
C(zipcode) [T.98011]    1.217e+05    1.45e+04     8.398     0.000     9.33e+04
1.5e+05
C(zipcode) [T.98014]    9.283e+04    1.71e+04     5.419     0.000     5.93e+04
1.26e+05
C(zipcode) [T.98019]    8.653e+04    1.46e+04     5.913     0.000     5.78e+04
1.15e+05
C(zipcode) [T.98022]   -5201.3953    1.38e+04    -0.376     0.707    -3.23e+04
2.19e+04
C(zipcode) [T.98023]   -3.287e+04    1.13e+04    -2.917     0.004    -5.5e+04
-1.08e+04
```

C(zipcode) [T.98024]	1.521e+05	2.03e+04	7.498	0.000	1.12e+05
1.92e+05					
C(zipcode) [T.98027]	1.708e+05	1.18e+04	14.465	0.000	1.48e+05
1.94e+05					
C(zipcode) [T.98028]	1.195e+05	1.29e+04	9.237	0.000	9.42e+04
1.45e+05					
C(zipcode) [T.98029]	2.105e+05	1.26e+04	16.720	0.000	1.86e+05
2.35e+05					
C(zipcode) [T.98030]	4456.9119	1.33e+04	0.335	0.738	-2.16e+04
3.05e+04					
C(zipcode) [T.98031]	1.854e+04	1.31e+04	1.419	0.156	-7068.667
4.42e+04					
C(zipcode) [T.98032]	1.034e+04	1.69e+04	0.611	0.541	-2.28e+04
4.35e+04					
C(zipcode) [T.98033]	3.705e+05	1.17e+04	31.779	0.000	3.48e+05
3.93e+05					
C(zipcode) [T.98034]	2.012e+05	1.11e+04	18.173	0.000	1.79e+05
2.23e+05					
C(zipcode) [T.98038]	3.168e+04	1.09e+04	2.900	0.004	1.03e+04
5.31e+04					
C(zipcode) [T.98039]	1.332e+06	2.48e+04	53.757	0.000	1.28e+06
1.38e+06					
C(zipcode) [T.98040]	5.297e+05	1.32e+04	40.275	0.000	5.04e+05
5.56e+05					
C(zipcode) [T.98042]	1.124e+04	1.1e+04	1.017	0.309	-1.04e+04
3.29e+04					
C(zipcode) [T.98045]	8.461e+04	1.4e+04	6.055	0.000	5.72e+04
1.12e+05					
C(zipcode) [T.98052]	2.285e+05	1.1e+04	20.786	0.000	2.07e+05
2.5e+05					
C(zipcode) [T.98053]	1.864e+05	1.19e+04	15.637	0.000	1.63e+05
2.1e+05					
C(zipcode) [T.98055]	4.632e+04	1.31e+04	3.524	0.000	2.06e+04
7.21e+04					
C(zipcode) [T.98056]	1.046e+05	1.18e+04	8.876	0.000	8.15e+04
1.28e+05					
C(zipcode) [T.98058]	3.19e+04	1.15e+04	2.777	0.005	9383.496
5.44e+04					
C(zipcode) [T.98059]	8.818e+04	1.14e+04	7.708	0.000	6.58e+04
1.11e+05					
C(zipcode) [T.98065]	8.139e+04	1.27e+04	6.406	0.000	5.65e+04
1.06e+05					
C(zipcode) [T.98070]	-1441.2949	1.77e+04	-0.082	0.935	-3.61e+04
3.32e+04					
C(zipcode) [T.98072]	1.516e+05	1.31e+04	11.569	0.000	1.26e+05
1.77e+05					
C(zipcode) [T.98074]	1.693e+05	1.17e+04	14.510	0.000	1.46e+05

1.92e+05					
C(zipcode) [T.98075]	1.696e+05	1.23e+04	13.793	0.000	1.46e+05
1.94e+05					
C(zipcode) [T.98077]	1.205e+05	1.46e+04	8.274	0.000	9.19e+04
1.49e+05					
C(zipcode) [T.98092]	-3.7e+04	1.22e+04	-3.025	0.002	-6.1e+04
-1.3e+04					
C(zipcode) [T.98102]	5.228e+05	1.85e+04	28.301	0.000	4.87e+05
5.59e+05					
C(zipcode) [T.98103]	3.504e+05	1.12e+04	31.350	0.000	3.29e+05
3.72e+05					
C(zipcode) [T.98105]	4.822e+05	1.41e+04	34.288	0.000	4.55e+05
5.1e+05					
C(zipcode) [T.98106]	1.485e+05	1.24e+04	11.958	0.000	1.24e+05
1.73e+05					
C(zipcode) [T.98107]	3.578e+05	1.34e+04	26.666	0.000	3.32e+05
3.84e+05					
C(zipcode) [T.98108]	1.291e+05	1.48e+04	8.729	0.000	1e+05
1.58e+05					
C(zipcode) [T.98109]	5.064e+05	1.81e+04	28.019	0.000	4.71e+05
5.42e+05					
C(zipcode) [T.98112]	6.362e+05	1.35e+04	46.985	0.000	6.1e+05
6.63e+05					
C(zipcode) [T.98115]	3.391e+05	1.11e+04	30.524	0.000	3.17e+05
3.61e+05					
C(zipcode) [T.98116]	3.02e+05	1.26e+04	23.964	0.000	2.77e+05
3.27e+05					
C(zipcode) [T.98117]	3.261e+05	1.12e+04	29.006	0.000	3.04e+05
3.48e+05					
C(zipcode) [T.98118]	1.763e+05	1.13e+04	15.536	0.000	1.54e+05
1.98e+05					
C(zipcode) [T.98119]	4.892e+05	1.51e+04	32.445	0.000	4.6e+05
5.19e+05					
C(zipcode) [T.98122]	3.503e+05	1.31e+04	26.696	0.000	3.25e+05
3.76e+05					
C(zipcode) [T.98125]	2.044e+05	1.18e+04	17.277	0.000	1.81e+05
2.28e+05					
C(zipcode) [T.98126]	1.996e+05	1.23e+04	16.212	0.000	1.75e+05
2.24e+05					
C(zipcode) [T.98133]	1.749e+05	1.13e+04	15.418	0.000	1.53e+05
1.97e+05					
C(zipcode) [T.98136]	2.535e+05	1.34e+04	18.985	0.000	2.27e+05
2.8e+05					
C(zipcode) [T.98144]	2.913e+05	1.25e+04	23.344	0.000	2.67e+05
3.16e+05					
C(zipcode) [T.98146]	1.137e+05	1.29e+04	8.803	0.000	8.84e+04
1.39e+05					

C(zipcode) [T.98148]	6.286e+04	2.32e+04	2.707	0.007	1.73e+04
1.08e+05					
C(zipcode) [T.98155]	1.512e+05	1.16e+04	13.071	0.000	1.29e+05
1.74e+05					
C(zipcode) [T.98166]	6.304e+04	1.34e+04	4.702	0.000	3.68e+04
8.93e+04					
C(zipcode) [T.98168]	7.529e+04	1.32e+04	5.710	0.000	4.94e+04
1.01e+05					
C(zipcode) [T.98177]	2.229e+05	1.34e+04	16.575	0.000	1.97e+05
2.49e+05					
C(zipcode) [T.98178]	3.638e+04	1.33e+04	2.736	0.006	1.03e+04
6.24e+04					
C(zipcode) [T.98188]	3.495e+04	1.64e+04	2.131	0.033	2804.530
6.71e+04					
C(zipcode) [T.98198]	1934.5695	1.3e+04	0.149	0.882	-2.36e+04
2.74e+04					
C(zipcode) [T.98199]	3.956e+05	1.28e+04	30.999	0.000	3.71e+05
4.21e+05					
bedrooms	-2.617e+04	1556.016	-16.822	0.000	-2.92e+04
-2.31e+04					
bathrooms	2.389e+04	2654.214	9.000	0.000	1.87e+04
2.91e+04					
sqft_lot	0.1804	0.030	6.111	0.000	0.123
0.238					
floors	-5.149e+04	3172.523	-16.230	0.000	-5.77e+04
-4.53e+04					
waterfront	6.922e+05	1.49e+04	46.571	0.000	6.63e+05
7.21e+05					
view	5.891e+04	1731.680	34.020	0.000	5.55e+04
6.23e+04					
grade	5.668e+04	1774.309	31.943	0.000	5.32e+04
6.02e+04					
sqft_above	213.8100	2.836	75.387	0.000	208.251
219.369					
sqft_basement	136.9940	3.516	38.962	0.000	130.102
143.886					
last_building_phase	-744.5707	59.173	-12.583	0.000	-860.555
-628.586					
=====					
Omnibus:	20673.566	Durbin-Watson:		1.992	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		4177058.103	
Skew:	4.117	Prob(JB):		0.00	
Kurtosis:	70.632	Cond. No.		4.54e+06	
=====					

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly

specified.

[2] The condition number is large, 4.54e+06. This might indicate that there are strong multicollinearity or other numerical problems.


"""

R-squared is 0.804 so we got an better result

So making dummies out of a categorical variable improved our model quite good. So lets take a look at other categorical variables.

Lets take a look at grade

```
[38]: sns.scatterplot(x="grade", y="price", data=df3);
```



Lets generate dummy variables for grade

```
[39]: model = ols('price ~ C(grade) + C(zipcode) + bedrooms + bathrooms + sqft_lot +  
    ↪ floors + waterfront + view + sqft_above + sqft_basement +  
    ↪ last_building_phase', data=df3).fit()  
  
model.summary()
```

```
[39]: <class 'statsmodels.iolib.summary.Summary'>  
      """
```

OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:          0.830
Model:                  OLS      Adj. R-squared:      0.829
Method:                 Least Squares    F-statistic:        1191.
Date:                   Sun, 03 Nov 2019    Prob (F-statistic):    0.00
Time:                   10:10:55    Log-Likelihood:       -2.8827e+05
No. Observations:      21597    AIC:                  5.767e+05
Df Residuals:          21508    BIC:                  5.774e+05
Df Model:               88
Covariance Type:       nonrobust
=====

```

```

=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
Intercept      9.453e+05    1.86e+05     5.088     0.000     5.81e+05
1.31e+06
C(grade) [T.4] -1.677e+05    1.55e+05    -1.081     0.280    -4.72e+05
1.36e+05
C(grade) [T.5] -2.039e+05    1.53e+05    -1.336     0.181    -5.03e+05
9.51e+04
C(grade) [T.6] -2.024e+05    1.52e+05    -1.329     0.184    -5.01e+05
9.61e+04
C(grade) [T.7] -2.007e+05    1.52e+05    -1.318     0.187    -4.99e+05
9.78e+04
C(grade) [T.8] -1.772e+05    1.52e+05    -1.163     0.245    -4.76e+05
1.21e+05
C(grade) [T.9] -1.031e+05    1.52e+05    -0.677     0.499    -4.02e+05
1.96e+05
C(grade) [T.10]  1.58e+04    1.52e+05     0.104     0.917    -2.83e+05
3.15e+05
C(grade) [T.11]  2.085e+05    1.53e+05     1.366     0.172    -9.08e+04
5.08e+05
C(grade) [T.12]  6.141e+05    1.54e+05     4.000     0.000     3.13e+05
9.15e+05
C(grade) [T.13]  1.712e+06    1.59e+05    10.785     0.000     1.4e+06
2.02e+06
C(zipcode) [T.98002]  1.377e+04    1.34e+04     1.024     0.306    -1.26e+04
4.01e+04
C(zipcode) [T.98003] -7461.9934    1.21e+04    -0.616     0.538    -3.12e+04
1.63e+04
C(zipcode) [T.98004]  7.787e+05    1.18e+04    65.755     0.000     7.55e+05
8.02e+05
C(zipcode) [T.98005]  3.198e+05    1.43e+04    22.328     0.000     2.92e+05
3.48e+05
C(zipcode) [T.98006]  2.635e+05    1.07e+04    24.660     0.000     2.43e+05

```

2.84e+05					
C(zipcode) [T.98007]	2.622e+05	1.51e+04	17.319	0.000	2.33e+05
2.92e+05					
C(zipcode) [T.98008]	2.678e+05	1.21e+04	22.082	0.000	2.44e+05
2.92e+05					
C(zipcode) [T.98010]	6.984e+04	1.72e+04	4.056	0.000	3.61e+04
1.04e+05					
C(zipcode) [T.98011]	1.372e+05	1.35e+04	10.148	0.000	1.11e+05
1.64e+05					
C(zipcode) [T.98014]	8.581e+04	1.6e+04	5.362	0.000	5.44e+04
1.17e+05					
C(zipcode) [T.98019]	9.398e+04	1.37e+04	6.882	0.000	6.72e+04
1.21e+05					
C(zipcode) [T.98022]	-2594.8172	1.29e+04	-0.201	0.841	-2.79e+04
2.27e+04					
C(zipcode) [T.98023]	-2.733e+04	1.05e+04	-2.598	0.009	-4.79e+04
-6708.254					
C(zipcode) [T.98024]	1.518e+05	1.89e+04	8.013	0.000	1.15e+05
1.89e+05					
C(zipcode) [T.98027]	1.732e+05	1.1e+04	15.701	0.000	1.52e+05
1.95e+05					
C(zipcode) [T.98028]	1.335e+05	1.21e+04	11.054	0.000	1.1e+05
1.57e+05					
C(zipcode) [T.98029]	2.239e+05	1.18e+04	19.034	0.000	2.01e+05
2.47e+05					
C(zipcode) [T.98030]	7077.1394	1.24e+04	0.570	0.569	-1.73e+04
3.14e+04					
C(zipcode) [T.98031]	2.154e+04	1.22e+04	1.767	0.077	-2357.074
4.54e+04					
C(zipcode) [T.98032]	8549.9058	1.58e+04	0.541	0.588	-2.24e+04
3.95e+04					
C(zipcode) [T.98033]	3.69e+05	1.09e+04	33.910	0.000	3.48e+05
3.9e+05					
C(zipcode) [T.98034]	2.018e+05	1.03e+04	19.536	0.000	1.82e+05
2.22e+05					
C(zipcode) [T.98038]	3.824e+04	1.02e+04	3.753	0.000	1.83e+04
5.82e+04					
C(zipcode) [T.98039]	1.245e+06	2.32e+04	53.722	0.000	1.2e+06
1.29e+06					
C(zipcode) [T.98040]	5.305e+05	1.23e+04	43.167	0.000	5.06e+05
5.55e+05					
C(zipcode) [T.98042]	1.398e+04	1.03e+04	1.357	0.175	-6220.325
3.42e+04					
C(zipcode) [T.98045]	8.979e+04	1.3e+04	6.888	0.000	6.42e+04
1.15e+05					
C(zipcode) [T.98052]	2.447e+05	1.03e+04	23.827	0.000	2.25e+05
2.65e+05					

C(zipcode) [T.98053] 2.28e+05	2.06e+05	1.11e+04	18.502	0.000	1.84e+05
C(zipcode) [T.98055] 6.59e+04	4.189e+04	1.23e+04	3.415	0.001	1.78e+04
C(zipcode) [T.98056] 1.23e+05	1.01e+05	1.1e+04	9.170	0.000	7.94e+04
C(zipcode) [T.98058] 5.96e+04	3.857e+04	1.07e+04	3.597	0.000	1.75e+04
C(zipcode) [T.98059] 1.1e+05	8.893e+04	1.07e+04	8.325	0.000	6.8e+04
C(zipcode) [T.98065] 1.16e+05	9.264e+04	1.19e+04	7.814	0.000	6.94e+04
C(zipcode) [T.98070] 4.17e+04	9401.1544	1.65e+04	0.570	0.569	-2.29e+04
C(zipcode) [T.98072] 1.89e+05	1.653e+05	1.22e+04	13.518	0.000	1.41e+05
C(zipcode) [T.98074] 2.02e+05	1.803e+05	1.09e+04	16.535	0.000	1.59e+05
C(zipcode) [T.98075] 2e+05	1.774e+05	1.15e+04	15.405	0.000	1.55e+05
C(zipcode) [T.98077] 1.51e+05	1.239e+05	1.36e+04	9.115	0.000	9.72e+04
C(zipcode) [T.98092] -335.283	-2.272e+04	1.14e+04	-1.989	0.047	-4.51e+04
C(zipcode) [T.98102] 5.33e+05	4.997e+05	1.73e+04	28.957	0.000	4.66e+05
C(zipcode) [T.98103] 3.74e+05	3.534e+05	1.04e+04	33.872	0.000	3.33e+05
C(zipcode) [T.98105] 5.2e+05	4.94e+05	1.31e+04	37.644	0.000	4.68e+05
C(zipcode) [T.98106] 1.46e+05	1.228e+05	1.16e+04	10.574	0.000	1e+05
C(zipcode) [T.98107] 3.81e+05	3.563e+05	1.25e+04	28.454	0.000	3.32e+05
C(zipcode) [T.98108] 1.46e+05	1.194e+05	1.38e+04	8.651	0.000	9.23e+04
C(zipcode) [T.98109] 5.51e+05	5.18e+05	1.69e+04	30.713	0.000	4.85e+05
C(zipcode) [T.98112] 6.74e+05	6.487e+05	1.26e+04	51.337	0.000	6.24e+05
C(zipcode) [T.98115] 3.66e+05	3.461e+05	1.04e+04	33.380	0.000	3.26e+05
C(zipcode) [T.98116] 3.3e+05	3.072e+05	1.18e+04	26.122	0.000	2.84e+05
C(zipcode) [T.98117] 3.48e+05	3.27e+05	1.05e+04	31.181	0.000	3.06e+05
C(zipcode) [T.98118]	1.635e+05	1.06e+04	15.414	0.000	1.43e+05

1.84e+05					
C(zipcode) [T.98119]	5.017e+05	1.41e+04	35.654	0.000	4.74e+05
5.29e+05					
C(zipcode) [T.98122]	3.588e+05	1.22e+04	29.301	0.000	3.35e+05
3.83e+05					
C(zipcode) [T.98125]	2.051e+05	1.1e+04	18.584	0.000	1.83e+05
2.27e+05					
C(zipcode) [T.98126]	1.915e+05	1.15e+04	16.650	0.000	1.69e+05
2.14e+05					
C(zipcode) [T.98133]	1.714e+05	1.06e+04	16.188	0.000	1.51e+05
1.92e+05					
C(zipcode) [T.98136]	2.576e+05	1.25e+04	20.680	0.000	2.33e+05
2.82e+05					
C(zipcode) [T.98144]	2.892e+05	1.16e+04	24.841	0.000	2.66e+05
3.12e+05					
C(zipcode) [T.98146]	9.82e+04	1.21e+04	8.129	0.000	7.45e+04
1.22e+05					
C(zipcode) [T.98148]	5.993e+04	2.17e+04	2.766	0.006	1.75e+04
1.02e+05					
C(zipcode) [T.98155]	1.5e+05	1.08e+04	13.887	0.000	1.29e+05
1.71e+05					
C(zipcode) [T.98166]	6.384e+04	1.25e+04	5.099	0.000	3.93e+04
8.84e+04					
C(zipcode) [T.98168]	4.702e+04	1.23e+04	3.808	0.000	2.28e+04
7.12e+04					
C(zipcode) [T.98177]	2.246e+05	1.26e+04	17.883	0.000	2e+05
2.49e+05					
C(zipcode) [T.98178]	2.344e+04	1.24e+04	1.887	0.059	-907.052
4.78e+04					
C(zipcode) [T.98188]	2.791e+04	1.53e+04	1.824	0.068	-2079.731
5.79e+04					
C(zipcode) [T.98198]	109.6074	1.21e+04	0.009	0.993	-2.37e+04
2.39e+04					
C(zipcode) [T.98199]	3.984e+05	1.19e+04	33.455	0.000	3.75e+05
4.22e+05					
bedrooms	-1.185e+04	1483.625	-7.989	0.000	-1.48e+04
-8944.530					
bathrooms	2.56e+04	2487.376	10.292	0.000	2.07e+04
3.05e+04					
sqft_lot	0.1719	0.028	6.240	0.000	0.118
0.226					
floors	-3.61e+04	2993.408	-12.061	0.000	-4.2e+04
-3.02e+04					
waterfront	6.73e+05	1.39e+04	48.441	0.000	6.46e+05
7e+05					
view	5.624e+04	1618.504	34.751	0.000	5.31e+04
5.94e+04					

sqft_above	173.3644	2.767	62.646	0.000	167.940
178.789					
sqft_basement	123.8584	3.290	37.651	0.000	117.410
130.306					
last_building_phase	-387.6541	55.856	-6.940	0.000	-497.136
-278.172					

```
=====
Omnibus:                16351.315    Durbin-Watson:                1.999
Prob(Omnibus):           0.000    Jarque-Bera (JB):            1841059.051
Skew:                    2.897    Prob(JB):                     0.00
Kurtosis:                47.859    Cond. No.                    2.16e+07
=====
```

Warnings:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.16e+07. This might indicate that there are strong multicollinearity or other numerical problems.

"""

Seems like a good model so far. Lets try to drop some variables that have a low correlation with the price. So we take a look at the heatmap again

```
[40]: corr = df2[['price', 'bedrooms', '
    ↳ 'zipcode', 'grade', 'bathrooms', 'sqft_lot', 'floors', 'waterfront', 'view', 'sqft_above', 'sqft_ba
    ↳ corr()
f, ax = plt.subplots(figsize = (18,18))
sns.heatmap(data=corr, center = 0, cmap="coolwarm", annot=True);
```



Lets try a model were we just keep variables with a correlation of 0.3 and above:

```
-bedrooms
-bathrooms
-sqft_above
-view
-sqft_basement
```

Also we want to keep our categorial variables

```
-zipcode
-grade
```

```
[41]: model = ols('price ~ C(grade) + C(zipcode) + bedrooms + bathrooms + sqft_above_
      ↪+ view + sqft_basement', data=df3).fit()

model.summary()
```

```
[41]: <class 'statsmodels.iolib.summary.Summary'>
      """
                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.809
Model:                            OLS    Adj. R-squared:              0.808
Method:                           Least Squares    F-statistic:                1084.
Date:                            Sun, 03 Nov 2019    Prob (F-statistic):          0.00
Time:                            10:10:56    Log-Likelihood:              -2.8952e+05
No. Observations:                 21597    AIC:                         5.792e+05
Df Residuals:                     21512    BIC:                         5.799e+05
Df Model:                          84
Covariance Type:                  nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept                1.809e+05    1.61e+05     1.120    0.263    -1.36e+05
4.97e+05
C(grade) [T.4]           -1.818e+05    1.64e+05    -1.107    0.268    -5.04e+05
1.4e+05
C(grade) [T.5]           -2.017e+05    1.62e+05    -1.248    0.212    -5.18e+05
1.15e+05
C(grade) [T.6]           -2.117e+05    1.61e+05    -1.313    0.189    -5.28e+05
1.04e+05
C(grade) [T.7]           -2.192e+05    1.61e+05    -1.359    0.174    -5.35e+05
9.69e+04
C(grade) [T.8]           -2.099e+05    1.61e+05    -1.301    0.193    -5.26e+05
1.06e+05
C(grade) [T.9]           -1.436e+05    1.61e+05    -0.890    0.374    -4.6e+05
1.73e+05
C(grade) [T.10]          -1.91e+04    1.61e+05    -0.118    0.906    -3.36e+05
2.97e+05
C(grade) [T.11]          1.816e+05    1.62e+05     1.123    0.261    -1.35e+05
4.99e+05
C(grade) [T.12]          6.193e+05    1.63e+05     3.810    0.000     3.01e+05
9.38e+05
C(grade) [T.13]          1.64e+06    1.68e+05     9.760    0.000     1.31e+06
1.97e+06
C(zipcode) [T.98002]     1.559e+04    1.42e+04     1.095    0.273    -1.23e+04
```


4.35e+04					
C(zipcode) [T.98003]	-4557.8566	1.28e+04	-0.355	0.722	-2.97e+04
2.06e+04					
C(zipcode) [T.98004]	7.85e+05	1.25e+04	62.754	0.000	7.61e+05
8.1e+05					
C(zipcode) [T.98005]	3.386e+05	1.51e+04	22.389	0.000	3.09e+05
3.68e+05					
C(zipcode) [T.98006]	2.609e+05	1.13e+04	23.099	0.000	2.39e+05
2.83e+05					
C(zipcode) [T.98007]	2.74e+05	1.6e+04	17.109	0.000	2.43e+05
3.05e+05					
C(zipcode) [T.98008]	2.916e+05	1.28e+04	22.759	0.000	2.66e+05
3.17e+05					
C(zipcode) [T.98010]	7.431e+04	1.82e+04	4.081	0.000	3.86e+04
1.1e+05					
C(zipcode) [T.98011]	1.412e+05	1.43e+04	9.863	0.000	1.13e+05
1.69e+05					
C(zipcode) [T.98014]	9.516e+04	1.68e+04	5.665	0.000	6.22e+04
1.28e+05					
C(zipcode) [T.98019]	9.333e+04	1.44e+04	6.462	0.000	6.5e+04
1.22e+05					
C(zipcode) [T.98022]	-2669.3218	1.36e+04	-0.197	0.844	-2.93e+04
2.39e+04					
C(zipcode) [T.98023]	-1.866e+04	1.11e+04	-1.675	0.094	-4.05e+04
3171.107					
C(zipcode) [T.98024]	1.693e+05	1.99e+04	8.500	0.000	1.3e+05
2.08e+05					
C(zipcode) [T.98027]	1.761e+05	1.17e+04	15.088	0.000	1.53e+05
1.99e+05					
C(zipcode) [T.98028]	1.372e+05	1.28e+04	10.731	0.000	1.12e+05
1.62e+05					
C(zipcode) [T.98029]	2.2e+05	1.24e+04	17.671	0.000	1.96e+05
2.44e+05					
C(zipcode) [T.98030]	9121.9831	1.32e+04	0.694	0.488	-1.67e+04
3.49e+04					
C(zipcode) [T.98031]	2.543e+04	1.29e+04	1.969	0.049	113.321
5.07e+04					
C(zipcode) [T.98032]	1.303e+04	1.67e+04	0.779	0.436	-1.97e+04
4.58e+04					
C(zipcode) [T.98033]	3.723e+05	1.15e+04	32.332	0.000	3.5e+05
3.95e+05					
C(zipcode) [T.98034]	2.1e+05	1.09e+04	19.206	0.000	1.89e+05
2.31e+05					
C(zipcode) [T.98038]	3.59e+04	1.08e+04	3.329	0.001	1.48e+04
5.7e+04					
C(zipcode) [T.98039]	1.265e+06	2.45e+04	51.592	0.000	1.22e+06
1.31e+06					

C(zipcode) [T.98040]	5.483e+05	1.3e+04	42.257	0.000	5.23e+05
5.74e+05					
C(zipcode) [T.98042]	1.603e+04	1.09e+04	1.468	0.142	-5368.883
3.74e+04					
C(zipcode) [T.98045]	9.131e+04	1.38e+04	6.629	0.000	6.43e+04
1.18e+05					
C(zipcode) [T.98052]	2.537e+05	1.09e+04	23.340	0.000	2.32e+05
2.75e+05					
C(zipcode) [T.98053]	2.115e+05	1.18e+04	17.969	0.000	1.88e+05
2.35e+05					
C(zipcode) [T.98055]	4.387e+04	1.3e+04	3.379	0.001	1.84e+04
6.93e+04					
C(zipcode) [T.98056]	1.039e+05	1.17e+04	8.907	0.000	8.1e+04
1.27e+05					
C(zipcode) [T.98058]	4.686e+04	1.14e+04	4.127	0.000	2.46e+04
6.91e+04					
C(zipcode) [T.98059]	9.082e+04	1.13e+04	8.026	0.000	6.86e+04
1.13e+05					
C(zipcode) [T.98065]	7.958e+04	1.26e+04	6.340	0.000	5.5e+04
1.04e+05					
C(zipcode) [T.98070]	1.119e+05	1.72e+04	6.498	0.000	7.81e+04
1.46e+05					
C(zipcode) [T.98072]	1.744e+05	1.29e+04	13.470	0.000	1.49e+05
2e+05					
C(zipcode) [T.98074]	1.91e+05	1.15e+04	16.546	0.000	1.68e+05
2.14e+05					
C(zipcode) [T.98075]	1.902e+05	1.22e+04	15.601	0.000	1.66e+05
2.14e+05					
C(zipcode) [T.98077]	1.391e+05	1.44e+04	9.685	0.000	1.11e+05
1.67e+05					
C(zipcode) [T.98092]	-1.979e+04	1.21e+04	-1.637	0.102	-4.35e+04
3900.651					
C(zipcode) [T.98102]	4.941e+05	1.8e+04	27.451	0.000	4.59e+05
5.29e+05					
C(zipcode) [T.98103]	3.443e+05	1.08e+04	32.031	0.000	3.23e+05
3.65e+05					
C(zipcode) [T.98105]	5.047e+05	1.36e+04	37.009	0.000	4.78e+05
5.31e+05					
C(zipcode) [T.98106]	1.188e+05	1.23e+04	9.677	0.000	9.47e+04
1.43e+05					
C(zipcode) [T.98107]	3.442e+05	1.31e+04	26.355	0.000	3.19e+05
3.7e+05					
C(zipcode) [T.98108]	1.195e+05	1.46e+04	8.209	0.000	9.1e+04
1.48e+05					
C(zipcode) [T.98109]	5.096e+05	1.76e+04	28.887	0.000	4.75e+05
5.44e+05					
C(zipcode) [T.98112]	6.514e+05	1.31e+04	49.860	0.000	6.26e+05

6.77e+05					
C(zipcode) [T.98115]	3.469e+05	1.08e+04	32.061	0.000	3.26e+05
3.68e+05					
C(zipcode) [T.98116]	2.94e+05	1.23e+04	23.846	0.000	2.7e+05
3.18e+05					
C(zipcode) [T.98117]	3.257e+05	1.09e+04	29.794	0.000	3.04e+05
3.47e+05					
C(zipcode) [T.98118]	1.637e+05	1.11e+04	14.685	0.000	1.42e+05
1.86e+05					
C(zipcode) [T.98119]	4.895e+05	1.47e+04	33.403	0.000	4.61e+05
5.18e+05					
C(zipcode) [T.98122]	3.536e+05	1.27e+04	27.744	0.000	3.29e+05
3.79e+05					
C(zipcode) [T.98125]	2.113e+05	1.16e+04	18.146	0.000	1.88e+05
2.34e+05					
C(zipcode) [T.98126]	1.837e+05	1.21e+04	15.171	0.000	1.6e+05
2.07e+05					
C(zipcode) [T.98133]	1.72e+05	1.12e+04	15.395	0.000	1.5e+05
1.94e+05					
C(zipcode) [T.98136]	2.545e+05	1.31e+04	19.407	0.000	2.29e+05
2.8e+05					
C(zipcode) [T.98144]	2.779e+05	1.22e+04	22.816	0.000	2.54e+05
3.02e+05					
C(zipcode) [T.98146]	1.061e+05	1.28e+04	8.303	0.000	8.1e+04
1.31e+05					
C(zipcode) [T.98148]	6.98e+04	2.29e+04	3.043	0.002	2.48e+04
1.15e+05					
C(zipcode) [T.98155]	1.619e+05	1.14e+04	14.188	0.000	1.4e+05
1.84e+05					
C(zipcode) [T.98166]	9.251e+04	1.32e+04	6.992	0.000	6.66e+04
1.18e+05					
C(zipcode) [T.98168]	4.958e+04	1.31e+04	3.797	0.000	2.4e+04
7.52e+04					
C(zipcode) [T.98177]	2.191e+05	1.33e+04	16.528	0.000	1.93e+05
2.45e+05					
C(zipcode) [T.98178]	3.658e+04	1.31e+04	2.785	0.005	1.08e+04
6.23e+04					
C(zipcode) [T.98188]	3.047e+04	1.62e+04	1.881	0.060	-1282.700
6.22e+04					
C(zipcode) [T.98198]	1.146e+04	1.29e+04	0.892	0.372	-1.37e+04
3.67e+04					
C(zipcode) [T.98199]	3.919e+05	1.25e+04	31.367	0.000	3.67e+05
4.16e+05					
bedrooms	-1.324e+04	1556.313	-8.505	0.000	-1.63e+04
-1.02e+04					
bathrooms	8127.4080	2351.084	3.457	0.001	3519.109
1.27e+04					

sqft_above	174.8861	2.826	61.887	0.000	169.347
180.425					
view	8.473e+04	1602.215	52.882	0.000	8.16e+04
8.79e+04					
sqft_basement	146.6707	3.205	45.765	0.000	140.389
152.952					

```

=====
Omnibus:                18948.630    Durbin-Watson:                2.001
Prob(Omnibus):           0.000    Jarque-Bera (JB):            2780571.536
Skew:                    3.629    Prob(JB):                     0.00
Kurtosis:                58.111    Cond. No.                     9.72e+05
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 9.72e+05. This might indicate that there are strong multicollinearity or other numerical problems.

"""

So we dropped five variabels and we still got a pretty nice R-squared value of 0.809