

# Immersive\_Data\_Science\_Project1\_EDA

November 15, 2019

## 0.1 Final Project Submission

Please fill out: \* Student name: Florian Fedrau

## 0.2 What to do

We are given the King County House Sales dataset and we want to predict the prices for houses.

## 0.3 Importing Python Libraries needed for the project

```
[1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
from pandas.plotting import scatter_matrix # data processing, scatter matrix
import statsmodels.api as sm # Python module that provides classes and
    ↪ functions for the estimation of many different
#statistical models, as well as for conducting statistical tests, and
    ↪ statistical data exploration
import matplotlib.pyplot as plt # plotting
import scipy.stats as stats
import sklearn.linear_model as linear_model
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold # import KFold
from sklearn.model_selection import cross_val_score, cross_val_predict
import seaborn as sns
import statsmodels.api as sms
import statsmodels.formula.api as smf
from statsmodels.formula.api import ols
from mpl_toolkits.basemap import Basemap
sns.set()
%matplotlib inline
```

## 0.4 Import the data

```
[2]: df = pd.read_csv('King_County_House_prices_dataset.csv')
df.head()
```

```
[2]:      id      date  price  bedrooms  bathrooms  sqft_living  \
0  7129300520  10/13/2014  221900.0         3         1.00        1180
```

1	6414100192	12/9/2014	538000.0	3	2.25	2570
2	5631500400	2/25/2015	180000.0	2	1.00	770
3	2487200875	12/9/2014	604000.0	4	3.00	1960
4	1954400510	2/18/2015	510000.0	3	2.00	1680

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
0	5650	1.0	NaN	0.0	...	7	1180	
1	7242	2.0	0.0	0.0	...	7	2170	
2	10000	1.0	0.0	0.0	...	6	770	
3	5000	1.0	0.0	0.0	...	7	1050	
4	8080	1.0	0.0	0.0	...	8	1680	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
0	0.0	1955	0.0	98178	47.5112	-122.257	
1	400.0	1951	1991.0	98125	47.7210	-122.319	
2	0.0	1933	NaN	98028	47.7379	-122.233	
3	910.0	1965	0.0	98136	47.5208	-122.393	
4	0.0	1987	0.0	98074	47.6168	-122.045	

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

[5 rows x 21 columns]

## 0.5 Understanding the data

```
[3]: f = open("column_names.md", "r")
      print(f.read())
      f.close()
```

```
# Column Names and descriptions for Kings County Data Set
* **id** - unique identified for a house
* **dateDate** - house was sold
* **pricePrice** - is prediction target
* **bedroomsNumber** - of Bedrooms/House
* **bathroomsNumber** - of bathrooms/bedrooms
* **sqft_livingsquare** - footage of the home
* **sqft_lotsquare** - footage of the lot
* **floorsTotal** - floors (levels) in house
* **waterfront** - House which has a view to a waterfront
* **view** - Has been viewed
* **condition** - How good the condition is ( Overall )
* **grade** - overall grade given to the housing unit, based on King County
```

```

grading system
* **sqft_above** - square footage of house apart from basement
* **sqft_basement** - square footage of the basement
* **yr_built** - Built Year
* **yr_renovated** - Year when house was renovated
* **zipcode** - zip
* **lat** - Latitude coordinate
* **long** - Longitude coordinate
* **sqft_living15** - The square footage of interior housing living space for
the nearest 15 neighbors
* **sqft_lot15** - The square footage of the land lots of the nearest 15
neighbors

```

## 0.6 Cleaning the data

Lets get an overview about our data first.

```
[4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
id                21597 non-null int64
date              21597 non-null object
price             21597 non-null float64
bedrooms          21597 non-null int64
bathrooms         21597 non-null float64
sqft_living       21597 non-null int64
sqft_lot          21597 non-null int64
floors            21597 non-null float64
waterfront        19221 non-null float64
view              21534 non-null float64
condition         21597 non-null int64
grade             21597 non-null int64
sqft_above        21597 non-null int64
sqft_basement     21597 non-null object
yr_built          21597 non-null int64
yr_renovated      17755 non-null float64
zipcode           21597 non-null int64
lat               21597 non-null float64
long              21597 non-null float64
sqft_living15     21597 non-null int64
sqft_lot15        21597 non-null int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB

```

Sort for nan values

```
[5]: df.isna().sum().sort_values()
```

```
[5]: id                0
    long              0
    lat              0
    zipcode          0
    yr_built         0
    sqft_basement    0
    sqft_above       0
    grade           0
    sqft_living15    0
    condition        0
    floors           0
    sqft_lot         0
    sqft_living      0
    bathrooms        0
    bedrooms         0
    price            0
    date             0
    sqft_lot15       0
    view             63
    waterfront       2376
    yr_renovated     3842
    dtype: int64
```

We see that there are missing values in the columns view, waterfront and yr\_renovated. Let's get an better understanding of the columns before we decide what to do with them.

## 0.7 Waterfront

First we take a look at the column waterfront

```
[6]: print("Unique values in waterfront:" + str(df.waterfront.unique()))
    print("Mean value of waterfront  :" + str(df.waterfront.mean()))
```

```
Unique values in waterfront:[nan  0.  1.]
Mean value of waterfront    :0.007595858696217679
```

Takeaway, basicly it is an boolean saying if a house has waterfront. We can also see that most of the houses have no waterfront.

We asume that a waterfront is an unique selling point and every houseowner would give us the information if the house has one. So we decide to fill the "nan" with zeros.

```
[7]: df.waterfront.fillna(0, inplace =True)
    df.waterfront.unique()
```

```
[7]: array([0., 1.])
```

## 0.8 View

We continue with the next column view

```
[8]: print("Unique values in view : " + str(df.view.unique()))
      print("Mean value of view : " + str(df.view.mean()))
      print("Number of Nan for view: " + str(df['view'].isna().sum()))
      print("Unique count for view : ")
      print(df.groupby('view')['id'].nunique())
```

```
Unique values in view : [ 0. nan  3.  4.  2.  1.]
Mean value of view : 0.23386272870808952
Number of Nan for view: 63
Unique count for view :
view
0.0    19253
1.0     329
2.0     956
3.0     505
4.0     314
Name: id, dtype: int64
```

We can see that most of the houses have not been viewed. We have only 63 nan in the column and we can fill them with zeros.

```
[9]: df.view.fillna(0, inplace =True)
      df.view.unique()
```

```
[9]: array([0., 3., 4., 2., 1.])
```

## 0.9 yr\_renovated

We continue with the column yr\_renovated

```
[10]: print("Unique values in yr_renovated : " + str(df.yr_renovated.unique()))
       print("Number of Nan for yr_renovated: " + str(df['yr_renovated'].isna().sum()))
       print("Unique count for yr_renovated : ")
       print(df.groupby('yr_renovated')['id'].nunique().head())
```

```
Unique values in yr_renovated : [  0. 1991.   nan 2002. 2010. 1992. 2013. 1994.
1978. 2005. 2003. 1984.
1954. 2014. 2011. 1983. 1945. 1990. 1988. 1977. 1981. 1995. 2000. 1999.
1998. 1970. 1989. 2004. 1986. 2007. 1987. 2006. 1985. 2001. 1980. 1971.
1979. 1997. 1950. 1969. 1948. 2009. 2015. 1974. 2008. 1968. 2012. 1963.
1951. 1962. 1953. 1993. 1996. 1955. 1982. 1956. 1940. 1976. 1946. 1975.
1964. 1973. 1957. 1959. 1960. 1967. 1965. 1934. 1972. 1944. 1958.]
Number of Nan for yr_renovated: 3842
Unique count for yr_renovated :
yr_renovated
0.0    16900
```

```
1934.0      1
1940.0      2
1944.0      1
1945.0      3
Name: id, dtype: int64
```

We assume that a renovation is a unique selling point and the homeowner would mention it. So we can fill the nan with a zero what means no renovation has taken place.

```
[11]: df.yr_renovated.fillna(0, inplace = True)
      df['yr_renovated'].isna().sum()
```

```
[11]: 0
```

So let's take a look at our data:

```
[12]: df.isna().sum().sort_values()
```

```
[12]: id          0
      long        0
      lat         0
      zipcode     0
      yr_renovated 0
      yr_built     0
      sqft_basement 0
      sqft_above   0
      grade        0
      sqft_living15 0
      condition    0
      waterfront   0
      floors       0
      sqft_lot     0
      sqft_living  0
      bathrooms    0
      bedrooms     0
      price        0
      date         0
      view         0
      sqft_lot15   0
      dtype: int64
```

We filled all of the missing data.

To be able to work with our data we want to have numeric values in our data. Let's figure out in what format our data is.

```
[13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
```

Data columns (total 21 columns):

id	21597	non-null	int64
date	21597	non-null	object
price	21597	non-null	float64
bedrooms	21597	non-null	int64
bathrooms	21597	non-null	float64
sqft_living	21597	non-null	int64
sqft_lot	21597	non-null	int64
floors	21597	non-null	float64
waterfront	21597	non-null	float64
view	21597	non-null	float64
condition	21597	non-null	int64
grade	21597	non-null	int64
sqft_above	21597	non-null	int64
sqft_basement	21597	non-null	object
yr_built	21597	non-null	int64
yr_renovated	21597	non-null	float64
zipcode	21597	non-null	int64
lat	21597	non-null	float64
long	21597	non-null	float64
sqft_living15	21597	non-null	int64
sqft_lot15	21597	non-null	int64

dtypes: float64(8), int64(11), object(2)  
memory usage: 3.5+ MB

## 0.10 sqft\_basement

So the columns date and sqft\_basement are objects. We want to change that. First we take a look at sqft\_basement. I would expect an numeric value. Lets explore sqft\_basement

```
[14]: print(df.groupby('sqft_basement')['id'].nunique())
```

sqft_basement	
0.0	12718
10.0	1
100.0	42
1000.0	146
1008.0	1
1010.0	62
1020.0	51
1024.0	1
1030.0	44
1040.0	54
1050.0	41
1060.0	58
1070.0	51
1080.0	31
1090.0	32

110.0	18
1100.0	78
1110.0	35
1120.0	43
1130.0	30
1135.0	1
1140.0	28
1150.0	26
1160.0	26
1170.0	30
1180.0	28
1190.0	24
120.0	53
1200.0	68
1210.0	18

...

80.0	20
800.0	201
810.0	55
820.0	62
830.0	56
840.0	82
850.0	69
860.0	79
861.0	1
862.0	1
870.0	48
875.0	1
880.0	69
890.0	52
90.0	21
900.0	141
906.0	1
910.0	69
915.0	1
920.0	65
930.0	41
935.0	1
940.0	71
946.0	1
950.0	62
960.0	65
970.0	44
980.0	55
990.0	51
?	454

Name: id, Length: 304, dtype: int64



We see the unique `sqft_basement` values and detect that in 454 rows is a “?”. That's why the datatype of the column is an object. We definitely want to change that.

So let's try to find some information about the houses with basement, without basement and we don't know about the basement.

```
[15]: df[['sqft_basement', 'sqft_above', 'sqft_living']].head()
```

```
[15]:   sqft_basement  sqft_above  sqft_living
0           0.0        1180        1180
1         400.0        2170        2570
2           0.0         770         770
3         910.0        1050        1960
4           0.0        1680        1680
```

```
[16]: basement = df[['sqft_basement', 'sqft_above', 'sqft_living']]
      basement['sqft_living-above'] = basement['sqft_living'] - basement['sqft_above']
      basement
```

```
/Users/flori/anaconda3/envs/nf/lib/python3.6/site-
packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
[16]:   sqft_basement  sqft_above  sqft_living  sqft_living-above
0           0.0        1180        1180             0
1         400.0        2170        2570          400
2           0.0         770         770             0
3         910.0        1050        1960          910
4           0.0        1680        1680             0
5        1530.0        3890        5420        1530
6            ?        1715        1715             0
7           0.0        1060        1060             0
8         730.0        1050        1780          730
9           0.0        1890        1890             0
10        1700.0        1860        3560        1700
11         300.0         860        1160          300
12           0.0        1430        1430             0
13           0.0        1370        1370             0
14           0.0        1810        1810             0
15         970.0        1980        2950          970
16           0.0        1890        1890             0
17           0.0        1600        1600             0
18            ?        1200        1200             0
```

19	0.0	1250	1250	0
20	760.0	860	1620	760
21	720.0	2330	3050	720
22	0.0	2270	2270	0
23	0.0	1070	1070	0
24	0.0	2450	2450	0
25	0.0	1710	1710	0
26	700.0	1750	2450	700
27	0.0	1400	1400	0
28	730.0	790	1520	730
29	0.0	2570	2570	0
...	...	...	...	...
21567	0.0	710	710	0
21568	320.0	940	1260	320
21569	0.0	1870	1870	0
21570	0.0	1430	1430	0
21571	0.0	1520	1520	0
21572	190.0	1020	1210	190
21573	0.0	2540	2540	0
21574	1800.0	3110	4910	1800
21575	0.0	2770	2770	0
21576	0.0	1190	1190	0
21577	0.0	4170	4170	0
21578	0.0	2500	2500	0
21579	50.0	1480	1530	50
21580	0.0	3600	3600	0
21581	?	3410	3410	0
21582	0.0	3118	3118	0
21583	0.0	3990	3990	0
21584	0.0	4470	4470	0
21585	0.0	1425	1425	0
21586	0.0	1500	1500	0
21587	0.0	2270	2270	0
21588	0.0	1490	1490	0
21589	0.0	2520	2520	0
21590	910.0	2600	3510	910
21591	130.0	1180	1310	130
21592	0.0	1530	1530	0
21593	0.0	2310	2310	0
21594	0.0	1020	1020	0
21595	0.0	1600	1600	0
21596	0.0	1020	1020	0

[21597 rows x 4 columns]

Seems like the sqft\_basement is just calculated by subtracting sqft\_above from sqft\_living

```
[17]: df['sqft_basement'] = basement['sqft_living'] - basement['sqft_above']
```

```
[18]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
id                21597 non-null int64
date              21597 non-null object
price             21597 non-null float64
bedrooms         21597 non-null int64
bathrooms        21597 non-null float64
sqft_living      21597 non-null int64
sqft_lot         21597 non-null int64
floors           21597 non-null float64
waterfront       21597 non-null float64
view             21597 non-null float64
condition        21597 non-null int64
grade            21597 non-null int64
sqft_above       21597 non-null int64
sqft_basement    21597 non-null int64
yr_built         21597 non-null int64
yr_renovated     21597 non-null float64
zipcode         21597 non-null int64
lat             21597 non-null float64
long            21597 non-null float64
sqft_living15    21597 non-null int64
sqft_lot15       21597 non-null int64
dtypes: float64(8), int64(12), object(1)
memory usage: 3.5+ MB
```

## 0.11 date

Now we have to look at the date column. To be able to work with it we have to change it to a numeric value. Therefore we want to separate between spring, summer, autumn and winter. We assume that the winter starts at december 1st, spring starts at march 1st, summer starts at june 1st and autumn at september 1st.

```
[19]: df['date'] = pd.to_datetime(df['date'])
```

```
[20]: dates = []
for i in range(len(df)):
    if(df['date'][i].month == 12 or df['date'][i].month == 1 or df['date'][i].
    month == 2):
        dates.append('winter')
    elif(df['date'][i].month == 3 or df['date'][i].month == 4 or df['date'][i].
    month == 5):
```

```

        dates.append('spring')
    elif(df['date'][i].month == 6 or df['date'][i].month == 7 or df['date'][i].
    ↪month == 8):
        dates.append('summer')
    else:
        dates.append('autumn')

```

```

[21]: dates_dummies = pd.get_dummies(dates, drop_first=True)
      dates_dummies.head()

```

```

[21]:   spring  summer  winter
0       0       0       0
1       0       0       1
2       0       0       1
3       0       0       1
4       0       0       1

```

We generated our dummy variables. Now we want to join our dummy variables with our dataset

```

[22]: df2 = pd.concat([df, dates_dummies], axis = 1)
      df2.head()

```

```

[22]:   id      date      price  bedrooms  bathrooms  sqft_living  \
0  7129300520  2014-10-13  221900.0         3         1.00         1180
1  6414100192  2014-12-09  538000.0         3         2.25         2570
2  5631500400  2015-02-25  180000.0         2         1.00          770
3  2487200875  2014-12-09  604000.0         4         3.00         1960
4  1954400510  2015-02-18  510000.0         3         2.00         1680

   sqft_lot  floors  waterfront  view  ...  yr_built  yr_renovated  \
0      5650     1.0         0.0  0.0  ...     1955         0.0
1      7242     2.0         0.0  0.0  ...     1951        1991.0
2     10000     1.0         0.0  0.0  ...     1933         0.0
3       5000     1.0         0.0  0.0  ...     1965         0.0
4       8080     1.0         0.0  0.0  ...     1987         0.0

   zipcode    lat    long  sqft_living15  sqft_lot15  spring  summer  \
0     98178  47.5112 -122.257         1340         5650       0       0
1     98125  47.7210 -122.319         1690         7639       0       0
2     98028  47.7379 -122.233         2720         8062       0       0
3     98136  47.5208 -122.393         1360         5000       0       0
4     98074  47.6168 -122.045         1800         7503       0       0

   winter
0       0
1       1
2       1

```

```
3      1
4      1
```

[5 rows x 24 columns]

Now we want to drop the month and day part of the date, so we just have the year left.

```
[23]: df2['date'] = pd.DatetimeIndex(df2['date']).year
```

```
[24]: df2.head()
```

```
[24]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	\
0	7129300520	2014	221900.0	3	1.00	1180	5650	
1	6414100192	2014	538000.0	3	2.25	2570	7242	
2	5631500400	2015	180000.0	2	1.00	770	10000	
3	2487200875	2014	604000.0	4	3.00	1960	5000	
4	1954400510	2015	510000.0	3	2.00	1680	8080	

	floors	waterfront	view	...	yr_built	yr_renovated	zipcode	lat	\
0	1.0	0.0	0.0	...	1955	0.0	98178	47.5112	
1	2.0	0.0	0.0	...	1951	1991.0	98125	47.7210	
2	1.0	0.0	0.0	...	1933	0.0	98028	47.7379	
3	1.0	0.0	0.0	...	1965	0.0	98136	47.5208	
4	1.0	0.0	0.0	...	1987	0.0	98074	47.6168	

	long	sqft_living15	sqft_lot15	spring	summer	winter
0	-122.257	1340	5650	0	0	0
1	-122.319	1690	7639	0	0	1
2	-122.233	2720	8062	0	0	1
3	-122.393	1360	5000	0	0	1
4	-122.045	1800	7503	0	0	1


[5 rows x 24 columns]

## 0.12 generate new feature: last\_building\_phase

Create new column “last\_building\_phase” include the max of yr\_built and yr\_renovated

```
[25]: df2["last_building_phase"] = df2[["yr_built", "yr_renovated"]].max(axis=1)
```

```
[26]: sns.scatterplot(x="last_building_phase", y="price", data=df2);
```




There is a small relationship between price and last\_building\_phase.

### 0.13 Visualisation

Lets plot some data

```
[27]: corr = df2.corr()
f, ax = plt.subplots(figsize = (18,18))
sns.heatmap(data=corr, center = 0, cmap="coolwarm", annot=True);
```



There are variables that have a correlation with price: - bedrooms 0.31 - bathrooms 0.53 - sqft\_living 0.7 - sqft\_lot 0.09 - floors 0.26 - waterfront 0.26 - view 0.39 - grade 0.67 - sqft\_above 0.31 - sqft\_basement 0.32 - year\_renovated 0.12 - lat 0.31 - sqft\_living15 0.59 - last\_building\_phase 0.098

Lets plot the location and the price:

```
[28]: # Extract the data we're interested in
lat = df['lat'].values
lon = df['long'].values
bins = [0, 150000, 200000, 250000, 300000, 400000, 550000, 750000, 1000000, 1500000, 2500000, 5000000, 10000000]
```


```

c = pd.cut(df['price'], bins, labels = range(12))
# 1. Draw the map background
fig = plt.figure(figsize=(16, 16))
m = Basemap(projection='merc', resolution='h', llcrnrlon = -122.6, llcrnrlat = 47.0,
urcrnrlon = -121.2, urcrnrlat = 47.8,
lon_0= -121.7,
lat_0= 47.4,
epsg=4269)
m.arcgisimage(service='ESRI_Imagery_World_2D', xpixels = 2000)

# 2. scatter house data, with color reflecting prizes
lons, lats = m(lon, lat)
m.scatter(lons, lats, zorder=1, linewidths=0.07, cmap = 'seismic', c = c)

plt.show()

```



We see that there is a correlation between location and price.

## 0.14 generate dummies for longitude and latitude

Lets group our houses for longitude and latitude

```

[29]: #seperate the latitude into seven parts
lat = []
for i in (range(len(df))):
    if df['lat'][i] < 47.2:

```



```

        lat.append('lat_47.2')
    elif df['lat'][i] < 47.3:
        lat.append('lat_47.3')
    elif df['lat'][i] < 47.4:
        lat.append('lat_47.4')
    elif df['lat'][i] < 47.5:
        lat.append('lat_47.5')
    elif df['lat'][i] < 47.6:
        lat.append('lat_47.6')
    elif df['lat'][i] < 47.7:
        lat.append('lat_47.7')
    else:
        lat.append('lat_47.8')

```

[30]: *#seperate the longitude into seven parts*

```

long = []
for i in (range(len(df))):
    if df['long'][i] < -122.4:
        long.append('long_-122.4')
    elif df['long'][i] < -122.2:
        long.append('long_-122.2')
    elif df['long'][i] < -122:
        long.append('long_-122')
    elif df['long'][i] < -121.8:
        long.append('long_-121.8')
    elif df['long'][i] < -121.6:
        long.append('long_-121.6')
    elif df['long'][i] < -121.4:
        long.append('long_-121.4')
    else:
        long.append('long_-121.2')

```

[31]: *#introduce dummy variables*

```

lat_dummies = pd.get_dummies(lat, drop_first=True)
long_dummies = pd.get_dummies(long, drop_first=True)
long_dummies.head()

```

```

[31]:
   long_-121.4  long_-121.6  long_-121.8  long_-122  long_-122.2  long_-122.4
0             0             0             0         0             1             0
1             0             0             0         0             1             0
2             0             0             0         0             1             0
3             0             0             0         0             1             0
4             0             0             0         1             0             0

```

```

[32]: lat_long = pd.concat([lat_dummies, long_dummies], axis = 1)
      lat_long.head()

```

```
[32]:   lat_47.3  lat_47.4  lat_47.5  lat_47.6  lat_47.7  lat_47.8  long_-121.4  \
0         0         0         0         1         0         0         0
1         0         0         0         0         0         1         0
2         0         0         0         0         0         1         0
3         0         0         0         1         0         0         0
4         0         0         0         0         1         0         0

      long_-121.6  long_-121.8  long_-122  long_-122.2  long_-122.4
0              0              0          0          1          0
1              0              0          0          1          0
2              0              0          0          1          0
3              0              0          0          1          0
4              0              0          1          0          0
```

```
[33]: # join our dataframe with the dummies
df3 = pd.concat([df2, lat_long], axis = 1)
```

## 0.15 Model a linear Regression

First we take the variables with the highest correlation with the price and our dummies for the location

```
[34]: X = df3[['bedrooms', 'bathrooms', 'sqft_living',
            'sqft_lot', 'floors', 'waterfront', 'view', 'grade',
            'sqft_above', 'sqft_basement', 'lat_47.8', 'lat_47.3', 'lat_47.4',
            'lat_47.5', 'lat_47.6', 'lat_47.7', 'long_-122.4', 'long_-121.4',
            'long_-121.6', 'long_-121.8', 'long_-122', 'long_-122.2']]
y = df3[['price']]
X = sms.add_constant(X)
model = sms.OLS(y, X)
results = model.fit()
results.summary()
```

```
[34]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                0.727
Model:                            OLS    Adj. R-squared:           0.727
Method:                 Least Squares    F-statistic:                2740.
Date:                  Mon, 04 Nov 2019    Prob (F-statistic):          0.00
Time:                  15:56:31    Log-Likelihood:             -2.9336e+05
No. Observations:                21597    AIC:                        5.868e+05
Df Residuals:                    21575    BIC:                        5.869e+05
Df Model:                        21
Covariance Type:                nonrobust
=====
```

=					
	coef	std err	t	P> t	[0.025
0.975]					
-----					
-					
const	-4.409e+05	6.26e+04	-7.039	0.000	-5.64e+05
-3.18e+05					
bedrooms	-2.451e+04	1806.001	-13.570	0.000	-2.8e+04
-2.1e+04					
bathrooms	1.454e+04	2936.901	4.952	0.000	8786.795
2.03e+04					
sqft_living	128.5872	2.005	64.125	0.000	124.657
132.518					
sqft_lot	0.1172	0.034	3.485	0.000	0.051
0.183					
floors	-7.445e+04	3338.798	-22.299	0.000	-8.1e+04
-6.79e+04					
waterfront	6.575e+05	1.73e+04	38.000	0.000	6.24e+05
6.91e+05					
view	5.832e+04	1999.164	29.171	0.000	5.44e+04
6.22e+04					
grade	6.656e+04	1941.184	34.291	0.000	6.28e+04
7.04e+04					
sqft_above	101.6745	2.035	49.965	0.000	97.686
105.663					
sqft_basement	26.9127	2.530	10.636	0.000	21.953
31.872					
lat_47.8	1.164e+05	2.17e+04	5.370	0.000	7.39e+04
1.59e+05					
lat_47.3	-5.139e+04	2.25e+04	-2.281	0.023	-9.56e+04
-7235.817					
lat_47.4	-4.009e+04	2.17e+04	-1.845	0.065	-8.27e+04
2489.904					
lat_47.5	1.744e+04	2.17e+04	0.802	0.422	-2.52e+04
6.01e+04					
lat_47.6	1.766e+05	2.16e+04	8.192	0.000	1.34e+05
2.19e+05					
lat_47.7	3.217e+05	2.16e+04	14.869	0.000	2.79e+05
3.64e+05					
long_-122.4	6.164e+04	5.9e+04	1.046	0.296	-5.39e+04
1.77e+05					
long_-121.4	-1.014e+05	9.74e+04	-1.041	0.298	-2.92e+05
8.96e+04					
long_-121.6	2.389e+04	5.95e+04	0.401	0.688	-9.28e+04
1.41e+05					
long_-121.8	-9.844e+04	5.83e+04	-1.689	0.091	-2.13e+05
1.58e+04					

```

long_-122      -2.936e+04   5.81e+04   -0.506      0.613   -1.43e+05
8.45e+04
long_-122.2     4.69e+04    5.8e+04    0.809      0.419   -6.68e+04
1.61e+05
=====
Omnibus:                18904.678   Durbin-Watson:                1.990
Prob(Omnibus):           0.000   Jarque-Bera (JB):            2171070.288
Skew:                    3.701   Prob(JB):                     0.00
Kurtosis:                51.558   Cond. No.                     7.81e+16
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 6.89e-21. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

"""

R-squared is 0.727 not bad at all.

## 0.16 Improving our approach

We want to introduce dummy variables for the zipcode, because the zipcode is an integer but is actual a categorial variable.

First we plot the zipcode and try to get some insights

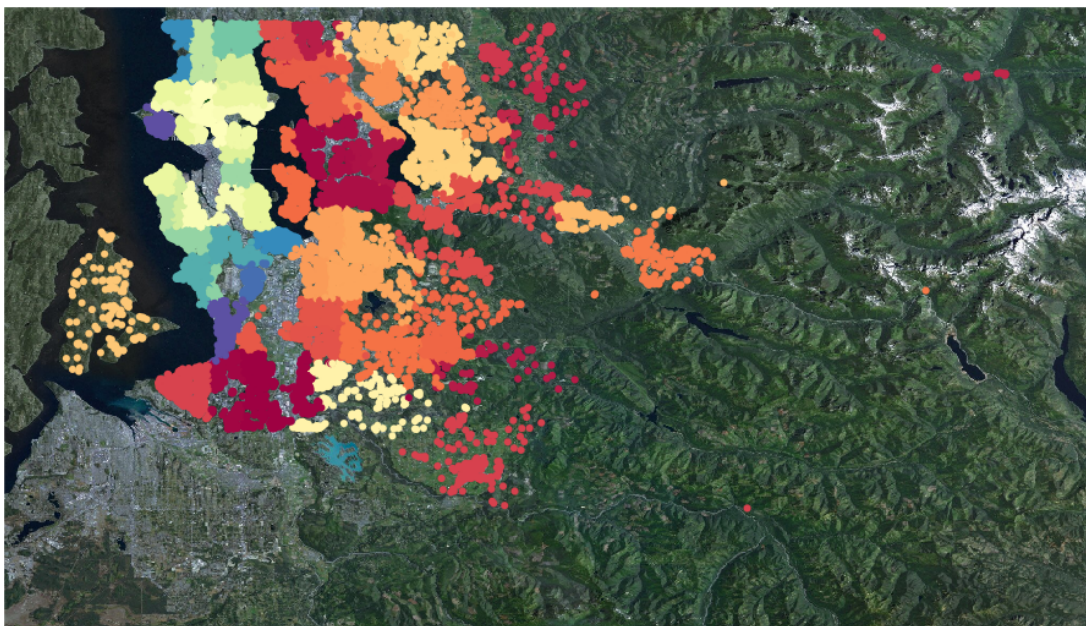
```

[35]: # Extract the data we're interested in
lat = df['lat'].values
lon = df['long'].values
# 1. Draw the map background
fig = plt.figure(figsize=(16, 16))
m = Basemap(projection='merc', resolution='h', llcrnrlon = -122.6, llcrnrlat = -47,
    urcrnrlon = -121.2, urcrnrlat = 47.8,
            lon_0= -121.7,
            lat_0= 47.4,
            epsg=4269)
m.arcgisimage(service='ESRI_Imagery_World_2D', xpixels = 2000)

# 2. scatter city data, with color reflecting zipcode areas
lons, lats = m(lon, lat)
m.scatter(lons, lats, zorder=1, linewidths=0.1, c = df['zipcode'].values, cmap=
    'Spectral')

plt.show()

```



Seems to work quite good and should be more exact than just the dummies for longitude and latitude.

```
[36]: #model with the most correlated variables and dummies for the zipcode
model = ols('price ~ C(zipcode) + bedrooms + bathrooms + sqft_lot + floors +
↳waterfront + view + grade + sqft_above + sqft_basement +
↳last_building_phase', data=df3).fit()

model.summary()
```

```
[36]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.804
Model:                            OLS    Adj. R-squared:              0.804
Method:                 Least Squares    F-statistic:                 1119.
Date:                  Mon, 04 Nov 2019    Prob (F-statistic):          0.00
Time:                  15:56:44    Log-Likelihood:             -2.8978e+05
No. Observations:          21597    AIC:                        5.797e+05
Df Residuals:              21517    BIC:                        5.804e+05
Df Model:                   79
Covariance Type:            nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
=====
```

0.975]

```
-----  
-----  
Intercept          1.056e+06  1.13e+05  9.346  0.000  8.34e+05  
1.28e+06  
C(zipcode) [T.98002] 3.497e+04  1.44e+04  2.429  0.015  6747.289  
6.32e+04  
C(zipcode) [T.98003] -1.719e+04  1.3e+04  -1.324  0.185  -4.26e+04  
8253.120  
C(zipcode) [T.98004] 7.868e+05  1.27e+04  62.025  0.000  7.62e+05  
8.12e+05  
C(zipcode) [T.98005] 3.104e+05  1.53e+04  20.244  0.000  2.8e+05  
3.41e+05  
C(zipcode) [T.98006] 2.772e+05  1.14e+04  24.239  0.000  2.55e+05  
3e+05  
C(zipcode) [T.98007] 2.513e+05  1.62e+04  15.488  0.000  2.19e+05  
2.83e+05  
C(zipcode) [T.98008] 2.538e+05  1.3e+04  19.529  0.000  2.28e+05  
2.79e+05  
C(zipcode) [T.98010] 7.7e+04  1.84e+04  4.174  0.000  4.08e+04  
1.13e+05  
C(zipcode) [T.98011] 1.217e+05  1.45e+04  8.398  0.000  9.33e+04  
1.5e+05  
C(zipcode) [T.98014] 9.283e+04  1.71e+04  5.419  0.000  5.93e+04  
1.26e+05  
C(zipcode) [T.98019] 8.653e+04  1.46e+04  5.913  0.000  5.78e+04  
1.15e+05  
C(zipcode) [T.98022] -5201.3953  1.38e+04  -0.376  0.707  -3.23e+04  
2.19e+04  
C(zipcode) [T.98023] -3.287e+04  1.13e+04  -2.917  0.004  -5.5e+04  
-1.08e+04  
C(zipcode) [T.98024] 1.521e+05  2.03e+04  7.498  0.000  1.12e+05  
1.92e+05  
C(zipcode) [T.98027] 1.708e+05  1.18e+04  14.465  0.000  1.48e+05  
1.94e+05  
C(zipcode) [T.98028] 1.195e+05  1.29e+04  9.237  0.000  9.42e+04  
1.45e+05  
C(zipcode) [T.98029] 2.105e+05  1.26e+04  16.720  0.000  1.86e+05  
2.35e+05  
C(zipcode) [T.98030] 4456.9119  1.33e+04  0.335  0.738  -2.16e+04  
3.05e+04  
C(zipcode) [T.98031] 1.854e+04  1.31e+04  1.419  0.156  -7068.667  
4.42e+04  
C(zipcode) [T.98032] 1.034e+04  1.69e+04  0.611  0.541  -2.28e+04  
4.35e+04  
C(zipcode) [T.98033] 3.705e+05  1.17e+04  31.779  0.000  3.48e+05  
3.93e+05
```

C(zipcode) [T.98034]	2.012e+05	1.11e+04	18.173	0.000	1.79e+05
2.23e+05					
C(zipcode) [T.98038]	3.168e+04	1.09e+04	2.900	0.004	1.03e+04
5.31e+04					
C(zipcode) [T.98039]	1.332e+06	2.48e+04	53.757	0.000	1.28e+06
1.38e+06					
C(zipcode) [T.98040]	5.297e+05	1.32e+04	40.275	0.000	5.04e+05
5.56e+05					
C(zipcode) [T.98042]	1.124e+04	1.1e+04	1.017	0.309	-1.04e+04
3.29e+04					
C(zipcode) [T.98045]	8.461e+04	1.4e+04	6.055	0.000	5.72e+04
1.12e+05					
C(zipcode) [T.98052]	2.285e+05	1.1e+04	20.786	0.000	2.07e+05
2.5e+05					
C(zipcode) [T.98053]	1.864e+05	1.19e+04	15.637	0.000	1.63e+05
2.1e+05					
C(zipcode) [T.98055]	4.632e+04	1.31e+04	3.524	0.000	2.06e+04
7.21e+04					
C(zipcode) [T.98056]	1.046e+05	1.18e+04	8.876	0.000	8.15e+04
1.28e+05					
C(zipcode) [T.98058]	3.19e+04	1.15e+04	2.777	0.005	9383.496
5.44e+04					
C(zipcode) [T.98059]	8.818e+04	1.14e+04	7.708	0.000	6.58e+04
1.11e+05					
C(zipcode) [T.98065]	8.139e+04	1.27e+04	6.406	0.000	5.65e+04
1.06e+05					
C(zipcode) [T.98070]	-1441.2949	1.77e+04	-0.082	0.935	-3.61e+04
3.32e+04					
C(zipcode) [T.98072]	1.516e+05	1.31e+04	11.569	0.000	1.26e+05
1.77e+05					
C(zipcode) [T.98074]	1.693e+05	1.17e+04	14.510	0.000	1.46e+05
1.92e+05					
C(zipcode) [T.98075]	1.696e+05	1.23e+04	13.793	0.000	1.46e+05
1.94e+05					
C(zipcode) [T.98077]	1.205e+05	1.46e+04	8.274	0.000	9.19e+04
1.49e+05					
C(zipcode) [T.98092]	-3.7e+04	1.22e+04	-3.025	0.002	-6.1e+04
-1.3e+04					
C(zipcode) [T.98102]	5.228e+05	1.85e+04	28.301	0.000	4.87e+05
5.59e+05					
C(zipcode) [T.98103]	3.504e+05	1.12e+04	31.350	0.000	3.29e+05
3.72e+05					
C(zipcode) [T.98105]	4.822e+05	1.41e+04	34.288	0.000	4.55e+05
5.1e+05					
C(zipcode) [T.98106]	1.485e+05	1.24e+04	11.958	0.000	1.24e+05
1.73e+05					
C(zipcode) [T.98107]	3.578e+05	1.34e+04	26.666	0.000	3.32e+05

3.84e+05						
C(zipcode) [T.98108]	1.291e+05	1.48e+04	8.729	0.000	1e+05	
1.58e+05						
C(zipcode) [T.98109]	5.064e+05	1.81e+04	28.019	0.000	4.71e+05	
5.42e+05						
C(zipcode) [T.98112]	6.362e+05	1.35e+04	46.985	0.000	6.1e+05	
6.63e+05						
C(zipcode) [T.98115]	3.391e+05	1.11e+04	30.524	0.000	3.17e+05	
3.61e+05						
C(zipcode) [T.98116]	3.02e+05	1.26e+04	23.964	0.000	2.77e+05	
3.27e+05						
C(zipcode) [T.98117]	3.261e+05	1.12e+04	29.006	0.000	3.04e+05	
3.48e+05						
C(zipcode) [T.98118]	1.763e+05	1.13e+04	15.536	0.000	1.54e+05	
1.98e+05						
C(zipcode) [T.98119]	4.892e+05	1.51e+04	32.445	0.000	4.6e+05	
5.19e+05						
C(zipcode) [T.98122]	3.503e+05	1.31e+04	26.696	0.000	3.25e+05	
3.76e+05						
C(zipcode) [T.98125]	2.044e+05	1.18e+04	17.277	0.000	1.81e+05	
2.28e+05						
C(zipcode) [T.98126]	1.996e+05	1.23e+04	16.212	0.000	1.75e+05	
2.24e+05						
C(zipcode) [T.98133]	1.749e+05	1.13e+04	15.418	0.000	1.53e+05	
1.97e+05						
C(zipcode) [T.98136]	2.535e+05	1.34e+04	18.985	0.000	2.27e+05	
2.8e+05						
C(zipcode) [T.98144]	2.913e+05	1.25e+04	23.344	0.000	2.67e+05	
3.16e+05						
C(zipcode) [T.98146]	1.137e+05	1.29e+04	8.803	0.000	8.84e+04	
1.39e+05						
C(zipcode) [T.98148]	6.286e+04	2.32e+04	2.707	0.007	1.73e+04	
1.08e+05						
C(zipcode) [T.98155]	1.512e+05	1.16e+04	13.071	0.000	1.29e+05	
1.74e+05						
C(zipcode) [T.98166]	6.304e+04	1.34e+04	4.702	0.000	3.68e+04	
8.93e+04						
C(zipcode) [T.98168]	7.529e+04	1.32e+04	5.710	0.000	4.94e+04	
1.01e+05						
C(zipcode) [T.98177]	2.229e+05	1.34e+04	16.575	0.000	1.97e+05	
2.49e+05						
C(zipcode) [T.98178]	3.638e+04	1.33e+04	2.736	0.006	1.03e+04	
6.24e+04						
C(zipcode) [T.98188]	3.495e+04	1.64e+04	2.131	0.033	2804.530	
6.71e+04						
C(zipcode) [T.98198]	1934.5695	1.3e+04	0.149	0.882	-2.36e+04	
2.74e+04						



C(zipcode) [T.98199]	3.956e+05	1.28e+04	30.999	0.000	3.71e+05
4.21e+05					
bedrooms	-2.617e+04	1556.016	-16.822	0.000	-2.92e+04
-2.31e+04					
bathrooms	2.389e+04	2654.214	9.000	0.000	1.87e+04
2.91e+04					
sqft_lot	0.1804	0.030	6.111	0.000	0.123
0.238					
floors	-5.149e+04	3172.523	-16.230	0.000	-5.77e+04
-4.53e+04					
waterfront	6.922e+05	1.49e+04	46.571	0.000	6.63e+05
7.21e+05					
view	5.891e+04	1731.680	34.020	0.000	5.55e+04
6.23e+04					
grade	5.668e+04	1774.309	31.943	0.000	5.32e+04
6.02e+04					
sqft_above	213.8100	2.836	75.387	0.000	208.251
219.369					
sqft_basement	136.9940	3.516	38.962	0.000	130.102
143.886					
last_building_phase	-744.5707	59.173	-12.583	0.000	-860.555
-628.586					

```

=====
Omnibus:                20673.566    Durbin-Watson:                1.992
Prob(Omnibus):           0.000    Jarque-Bera (JB):            4177058.103
Skew:                    4.117    Prob(JB):                     0.00
Kurtosis:                70.632    Cond. No.                    4.54e+06
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.54e+06. This might indicate that there are strong multicollinearity or other numerical problems.


"""

R-squared is 0.804 so we got an better result

So making dummies out of a categorial variable improved our model quite good. So lets take a look at other categorial variables.

Lets take a look at grade

```
[37]: sns.scatterplot(x="grade", y="price", data=df3);
```



Lets generate dummy variables for grade and model it

```
[38]: model = ols('price ~ C(grade) + C(zipcode) + bedrooms + bathrooms+ sqft_living_
↳+ sqft_lot + floors + waterfront + view + sqft_above + sqft_basement +_
↳last_building_phase', data=df3).fit()
```

```
model.summary()
```

```
[38]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

#### OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.830
Model:                  OLS      Adj. R-squared:           0.829
Method:                 Least Squares    F-statistic:            1191.
Date:                   Mon, 04 Nov 2019    Prob (F-statistic):      0.00
Time:                   15:56:45    Log-Likelihood:         -2.8827e+05
No. Observations:       21597    AIC:                    5.767e+05
Df Residuals:           21508    BIC:                    5.774e+05
Df Model:                88
Covariance Type:        nonrobust
```

```
=====
=====
coef    std err          t    P>|t|    [0.025
```

0.975]

```
-----
-----
Intercept          9.453e+05  1.86e+05  5.088  0.000  5.81e+05
1.31e+06
C(grade) [T.4]     -1.677e+05  1.55e+05 -1.081  0.280 -4.72e+05
1.36e+05
C(grade) [T.5]     -2.039e+05  1.53e+05 -1.336  0.181 -5.03e+05
9.51e+04
C(grade) [T.6]     -2.024e+05  1.52e+05 -1.329  0.184 -5.01e+05
9.61e+04
C(grade) [T.7]     -2.007e+05  1.52e+05 -1.318  0.187 -4.99e+05
9.78e+04
C(grade) [T.8]     -1.772e+05  1.52e+05 -1.163  0.245 -4.76e+05
1.21e+05
C(grade) [T.9]     -1.031e+05  1.52e+05 -0.677  0.499 -4.02e+05
1.96e+05
C(grade) [T.10]    1.58e+04  1.52e+05  0.104  0.917 -2.83e+05
3.15e+05
C(grade) [T.11]    2.085e+05  1.53e+05  1.366  0.172 -9.08e+04
5.08e+05
C(grade) [T.12]    6.141e+05  1.54e+05  4.000  0.000  3.13e+05
9.15e+05
C(grade) [T.13]    1.712e+06  1.59e+05 10.785  0.000  1.4e+06
2.02e+06
C(zipcode) [T.98002] 1.377e+04  1.34e+04  1.024  0.306 -1.26e+04
4.01e+04
C(zipcode) [T.98003] -7461.9934  1.21e+04 -0.616  0.538 -3.12e+04
1.63e+04
C(zipcode) [T.98004] 7.787e+05  1.18e+04 65.755  0.000  7.55e+05
8.02e+05
C(zipcode) [T.98005] 3.198e+05  1.43e+04 22.328  0.000  2.92e+05
3.48e+05
C(zipcode) [T.98006] 2.635e+05  1.07e+04 24.660  0.000  2.43e+05
2.84e+05
C(zipcode) [T.98007] 2.622e+05  1.51e+04 17.319  0.000  2.33e+05
2.92e+05
C(zipcode) [T.98008] 2.678e+05  1.21e+04 22.082  0.000  2.44e+05
2.92e+05
C(zipcode) [T.98010] 6.984e+04  1.72e+04  4.056  0.000  3.61e+04
1.04e+05
C(zipcode) [T.98011] 1.372e+05  1.35e+04 10.148  0.000  1.11e+05
1.64e+05
C(zipcode) [T.98014] 8.581e+04  1.6e+04  5.362  0.000  5.44e+04
1.17e+05
C(zipcode) [T.98019] 9.398e+04  1.37e+04  6.882  0.000  6.72e+04
1.21e+05
```

C(zipcode) [T.98022]	-2594.8172	1.29e+04	-0.201	0.841	-2.79e+04
2.27e+04					
C(zipcode) [T.98023]	-2.733e+04	1.05e+04	-2.598	0.009	-4.79e+04
-6708.254					
C(zipcode) [T.98024]	1.518e+05	1.89e+04	8.013	0.000	1.15e+05
1.89e+05					
C(zipcode) [T.98027]	1.732e+05	1.1e+04	15.701	0.000	1.52e+05
1.95e+05					
C(zipcode) [T.98028]	1.335e+05	1.21e+04	11.054	0.000	1.1e+05
1.57e+05					
C(zipcode) [T.98029]	2.239e+05	1.18e+04	19.034	0.000	2.01e+05
2.47e+05					
C(zipcode) [T.98030]	7077.1394	1.24e+04	0.570	0.569	-1.73e+04
3.14e+04					
C(zipcode) [T.98031]	2.154e+04	1.22e+04	1.767	0.077	-2357.074
4.54e+04					
C(zipcode) [T.98032]	8549.9058	1.58e+04	0.541	0.588	-2.24e+04
3.95e+04					
C(zipcode) [T.98033]	3.69e+05	1.09e+04	33.910	0.000	3.48e+05
3.9e+05					
C(zipcode) [T.98034]	2.018e+05	1.03e+04	19.536	0.000	1.82e+05
2.22e+05					
C(zipcode) [T.98038]	3.824e+04	1.02e+04	3.753	0.000	1.83e+04
5.82e+04					
C(zipcode) [T.98039]	1.245e+06	2.32e+04	53.722	0.000	1.2e+06
1.29e+06					
C(zipcode) [T.98040]	5.305e+05	1.23e+04	43.167	0.000	5.06e+05
5.55e+05					
C(zipcode) [T.98042]	1.398e+04	1.03e+04	1.357	0.175	-6220.325
3.42e+04					
C(zipcode) [T.98045]	8.979e+04	1.3e+04	6.888	0.000	6.42e+04
1.15e+05					
C(zipcode) [T.98052]	2.447e+05	1.03e+04	23.827	0.000	2.25e+05
2.65e+05					
C(zipcode) [T.98053]	2.06e+05	1.11e+04	18.502	0.000	1.84e+05
2.28e+05					
C(zipcode) [T.98055]	4.189e+04	1.23e+04	3.415	0.001	1.78e+04
6.59e+04					
C(zipcode) [T.98056]	1.01e+05	1.1e+04	9.170	0.000	7.94e+04
1.23e+05					
C(zipcode) [T.98058]	3.857e+04	1.07e+04	3.597	0.000	1.75e+04
5.96e+04					
C(zipcode) [T.98059]	8.893e+04	1.07e+04	8.325	0.000	6.8e+04
1.1e+05					
C(zipcode) [T.98065]	9.264e+04	1.19e+04	7.814	0.000	6.94e+04
1.16e+05					
C(zipcode) [T.98070]	9401.1544	1.65e+04	0.570	0.569	-2.29e+04

4.17e+04						
C(zipcode) [T.98072]	1.653e+05	1.22e+04	13.518	0.000	1.41e+05	
1.89e+05						
C(zipcode) [T.98074]	1.803e+05	1.09e+04	16.535	0.000	1.59e+05	
2.02e+05						
C(zipcode) [T.98075]	1.774e+05	1.15e+04	15.405	0.000	1.55e+05	
2e+05						
C(zipcode) [T.98077]	1.239e+05	1.36e+04	9.115	0.000	9.72e+04	
1.51e+05						
C(zipcode) [T.98092]	-2.272e+04	1.14e+04	-1.989	0.047	-4.51e+04	
-335.283						
C(zipcode) [T.98102]	4.997e+05	1.73e+04	28.957	0.000	4.66e+05	
5.33e+05						
C(zipcode) [T.98103]	3.534e+05	1.04e+04	33.872	0.000	3.33e+05	
3.74e+05						
C(zipcode) [T.98105]	4.94e+05	1.31e+04	37.644	0.000	4.68e+05	
5.2e+05						
C(zipcode) [T.98106]	1.228e+05	1.16e+04	10.574	0.000	1e+05	
1.46e+05						
C(zipcode) [T.98107]	3.563e+05	1.25e+04	28.454	0.000	3.32e+05	
3.81e+05						
C(zipcode) [T.98108]	1.194e+05	1.38e+04	8.651	0.000	9.23e+04	
1.46e+05						
C(zipcode) [T.98109]	5.18e+05	1.69e+04	30.713	0.000	4.85e+05	
5.51e+05						
C(zipcode) [T.98112]	6.487e+05	1.26e+04	51.337	0.000	6.24e+05	
6.74e+05						
C(zipcode) [T.98115]	3.461e+05	1.04e+04	33.380	0.000	3.26e+05	
3.66e+05						
C(zipcode) [T.98116]	3.072e+05	1.18e+04	26.122	0.000	2.84e+05	
3.3e+05						
C(zipcode) [T.98117]	3.27e+05	1.05e+04	31.181	0.000	3.06e+05	
3.48e+05						
C(zipcode) [T.98118]	1.635e+05	1.06e+04	15.414	0.000	1.43e+05	
1.84e+05						
C(zipcode) [T.98119]	5.017e+05	1.41e+04	35.654	0.000	4.74e+05	
5.29e+05						
C(zipcode) [T.98122]	3.588e+05	1.22e+04	29.301	0.000	3.35e+05	
3.83e+05						
C(zipcode) [T.98125]	2.051e+05	1.1e+04	18.584	0.000	1.83e+05	
2.27e+05						
C(zipcode) [T.98126]	1.915e+05	1.15e+04	16.650	0.000	1.69e+05	
2.14e+05						
C(zipcode) [T.98133]	1.714e+05	1.06e+04	16.188	0.000	1.51e+05	
1.92e+05						
C(zipcode) [T.98136]	2.576e+05	1.25e+04	20.680	0.000	2.33e+05	
2.82e+05						

C(zipcode) [T.98144]	2.892e+05	1.16e+04	24.841	0.000	2.66e+05
3.12e+05					
C(zipcode) [T.98146]	9.82e+04	1.21e+04	8.129	0.000	7.45e+04
1.22e+05					
C(zipcode) [T.98148]	5.993e+04	2.17e+04	2.766	0.006	1.75e+04
1.02e+05					
C(zipcode) [T.98155]	1.5e+05	1.08e+04	13.887	0.000	1.29e+05
1.71e+05					
C(zipcode) [T.98166]	6.384e+04	1.25e+04	5.099	0.000	3.93e+04
8.84e+04					
C(zipcode) [T.98168]	4.702e+04	1.23e+04	3.808	0.000	2.28e+04
7.12e+04					
C(zipcode) [T.98177]	2.246e+05	1.26e+04	17.883	0.000	2e+05
2.49e+05					
C(zipcode) [T.98178]	2.344e+04	1.24e+04	1.887	0.059	-907.052
4.78e+04					
C(zipcode) [T.98188]	2.791e+04	1.53e+04	1.824	0.068	-2079.731
5.79e+04					
C(zipcode) [T.98198]	109.6074	1.21e+04	0.009	0.993	-2.37e+04
2.39e+04					
C(zipcode) [T.98199]	3.984e+05	1.19e+04	33.455	0.000	3.75e+05
4.22e+05					
bedrooms	-1.185e+04	1483.625	-7.989	0.000	-1.48e+04
-8944.530					
bathrooms	2.56e+04	2487.376	10.292	0.000	2.07e+04
3.05e+04					
sqft_living	99.0743	1.668	59.402	0.000	95.805
102.343					
sqft_lot	0.1719	0.028	6.240	0.000	0.118
0.226					
floors	-3.61e+04	2993.408	-12.061	0.000	-4.2e+04
-3.02e+04					
waterfront	6.73e+05	1.39e+04	48.441	0.000	6.46e+05
7e+05					
view	5.624e+04	1618.504	34.751	0.000	5.31e+04
5.94e+04					
sqft_above	74.2902	1.775	41.863	0.000	70.812
77.769					
sqft_basement	24.7841	2.050	12.088	0.000	20.765
28.803					
last_building_phase	-387.6541	55.856	-6.940	0.000	-497.136
-278.172					
=====					
Omnibus:	16351.315	Durbin-Watson:		1.999	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		1841059.051	
Skew:	2.897	Prob(JB):		0.00	
Kurtosis:	47.859	Cond. No.		2.25e+15	

=====


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 8.28e-18. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.  
"""

Seems like a good model so far. Lets try to drop some variables that have a low correlation with the price. So we take a look at the heatmap again

```
[39]: corr = df2[['price', 'bedrooms', 'zipcode', 'grade', 'bathrooms', 'sqft_living',  
    ↪ 'sqft_lot', 'floors', 'waterfront', 'view', 'sqft_above', 'sqft_basement', 'last_building_phase']  
    ↪ corr()  
f, ax = plt.subplots(figsize = (18,18))  
sns.heatmap(data=corr, center = 0, cmap="coolwarm", annot=True);
```



Lets try a model were we just keep variables with a correlation of 0.3 and above: - bathrooms - bedrooms - sqft\_living - sqft\_above - view - sqft\_basement

Also we want to keep our categorical variables - zipcode - grade

We also decided to drop bedrooms because bathrooms is defined as bathrooms per bedrooms therefore they have a high correlation. The same goes for sqft\_above and sqft\_basement, because the sum of them is sqft\_living



```
[40]: model = ols('price ~ C(grade) + C(zipcode) + bathrooms + sqft_living + view',
↳data=df3).fit()

model.summary()
```

```
[40]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.808
Model:                            OLS    Adj. R-squared:             0.807
Method:                 Least Squares    F-statistic:                 1101.
Date:                  Mon, 04 Nov 2019    Prob (F-statistic):          0.00
Time:                  15:56:46    Log-Likelihood:             -2.8959e+05
No. Observations:          21597    AIC:                        5.794e+05
Df Residuals:              21514    BIC:                        5.800e+05
Df Model:                   82
Covariance Type:            nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept                1.758e+05    1.62e+05     1.085    0.278    -1.42e+05
4.93e+05
C(grade) [T.4]          -1.812e+05    1.65e+05    -1.099    0.272    -5.04e+05
1.42e+05
C(grade) [T.5]          -2.053e+05    1.62e+05    -1.266    0.205    -5.23e+05
1.13e+05
C(grade) [T.6]          -2.155e+05    1.62e+05    -1.331    0.183    -5.33e+05
1.02e+05
C(grade) [T.7]          -2.239e+05    1.62e+05    -1.383    0.167    -5.41e+05
9.33e+04
C(grade) [T.8]          -2.074e+05    1.62e+05    -1.281    0.200    -5.25e+05
1.1e+05
C(grade) [T.9]          -1.307e+05    1.62e+05    -0.807    0.419    -4.48e+05
1.87e+05
C(grade) [T.10]         3924.9530    1.62e+05     0.024    0.981    -3.14e+05
3.21e+05
C(grade) [T.11]         2.168e+05    1.62e+05     1.337    0.181    -1.01e+05
5.35e+05
C(grade) [T.12]         6.706e+05    1.63e+05     4.113    0.000     3.51e+05
9.9e+05
C(grade) [T.13]         1.711e+06    1.69e+05    10.154    0.000     1.38e+06
2.04e+06
C(zipcode) [T.98002]    1.522e+04    1.43e+04     1.065    0.287    -1.28e+04

```

4.32e+04					
C(zipcode) [T.98003]	-6813.3564	1.29e+04	-0.529	0.597	-3.2e+04
1.84e+04					
C(zipcode) [T.98004]	7.798e+05	1.25e+04	62.174	0.000	7.55e+05
8.04e+05					
C(zipcode) [T.98005]	3.301e+05	1.52e+04	21.783	0.000	3e+05
3.6e+05					
C(zipcode) [T.98006]	2.524e+05	1.13e+04	22.332	0.000	2.3e+05
2.75e+05					
C(zipcode) [T.98007]	2.664e+05	1.61e+04	16.589	0.000	2.35e+05
2.98e+05					
C(zipcode) [T.98008]	2.835e+05	1.28e+04	22.076	0.000	2.58e+05
3.09e+05					
C(zipcode) [T.98010]	7.861e+04	1.83e+04	4.303	0.000	4.28e+04
1.14e+05					
C(zipcode) [T.98011]	1.403e+05	1.44e+04	9.769	0.000	1.12e+05
1.69e+05					
C(zipcode) [T.98014]	1.029e+05	1.68e+04	6.109	0.000	6.99e+04
1.36e+05					
C(zipcode) [T.98019]	9.942e+04	1.45e+04	6.864	0.000	7.1e+04
1.28e+05					
C(zipcode) [T.98022]	1622.5155	1.36e+04	0.119	0.905	-2.51e+04
2.83e+04					
C(zipcode) [T.98023]	-2.225e+04	1.12e+04	-1.992	0.046	-4.41e+04
-351.073					
C(zipcode) [T.98024]	1.757e+05	2e+04	8.797	0.000	1.37e+05
2.15e+05					
C(zipcode) [T.98027]	1.74e+05	1.17e+04	14.874	0.000	1.51e+05
1.97e+05					
C(zipcode) [T.98028]	1.346e+05	1.28e+04	10.495	0.000	1.09e+05
1.6e+05					
C(zipcode) [T.98029]	2.225e+05	1.25e+04	17.815	0.000	1.98e+05
2.47e+05					
C(zipcode) [T.98030]	9220.0878	1.32e+04	0.699	0.485	-1.66e+04
3.51e+04					
C(zipcode) [T.98031]	2.378e+04	1.3e+04	1.835	0.066	-1618.678
4.92e+04					
C(zipcode) [T.98032]	6383.0323	1.68e+04	0.381	0.703	-2.65e+04
3.93e+04					
C(zipcode) [T.98033]	3.7e+05	1.16e+04	32.031	0.000	3.47e+05
3.93e+05					
C(zipcode) [T.98034]	2.06e+05	1.1e+04	18.787	0.000	1.85e+05
2.27e+05					
C(zipcode) [T.98038]	4.166e+04	1.08e+04	3.854	0.000	2.05e+04
6.29e+04					
C(zipcode) [T.98039]	1.267e+06	2.46e+04	51.485	0.000	1.22e+06
1.31e+06					

C(zipcode) [T.98040]	5.395e+05	1.3e+04	41.518	0.000	5.14e+05
5.65e+05					
C(zipcode) [T.98042]	1.801e+04	1.1e+04	1.645	0.100	-3455.576
3.95e+04					
C(zipcode) [T.98045]	9.586e+04	1.38e+04	6.938	0.000	6.88e+04
1.23e+05					
C(zipcode) [T.98052]	2.519e+05	1.09e+04	23.099	0.000	2.31e+05
2.73e+05					
C(zipcode) [T.98053]	2.243e+05	1.18e+04	19.059	0.000	2.01e+05
2.47e+05					
C(zipcode) [T.98055]	4.32e+04	1.3e+04	3.316	0.001	1.77e+04
6.87e+04					
C(zipcode) [T.98056]	1.026e+05	1.17e+04	8.768	0.000	7.97e+04
1.26e+05					
C(zipcode) [T.98058]	4.434e+04	1.14e+04	3.892	0.000	2.2e+04
6.67e+04					
C(zipcode) [T.98059]	9.364e+04	1.13e+04	8.252	0.000	7.14e+04
1.16e+05					
C(zipcode) [T.98065]	8.97e+04	1.26e+04	7.138	0.000	6.51e+04
1.14e+05					
C(zipcode) [T.98070]	1.186e+05	1.73e+04	6.869	0.000	8.47e+04
1.52e+05					
C(zipcode) [T.98072]	1.745e+05	1.3e+04	13.436	0.000	1.49e+05
2e+05					
C(zipcode) [T.98074]	1.917e+05	1.16e+04	16.546	0.000	1.69e+05
2.14e+05					
C(zipcode) [T.98075]	1.938e+05	1.22e+04	15.850	0.000	1.7e+05
2.18e+05					
C(zipcode) [T.98077]	1.433e+05	1.44e+04	9.947	0.000	1.15e+05
1.72e+05					
C(zipcode) [T.98092]	-1.729e+04	1.21e+04	-1.425	0.154	-4.11e+04
6486.229					
C(zipcode) [T.98102]	4.874e+05	1.8e+04	27.022	0.000	4.52e+05
5.23e+05					
C(zipcode) [T.98103]	3.414e+05	1.08e+04	31.691	0.000	3.2e+05
3.63e+05					
C(zipcode) [T.98105]	4.958e+05	1.37e+04	36.305	0.000	4.69e+05
5.23e+05					
C(zipcode) [T.98106]	1.117e+05	1.23e+04	9.085	0.000	8.76e+04
1.36e+05					
C(zipcode) [T.98107]	3.405e+05	1.31e+04	26.014	0.000	3.15e+05
3.66e+05					
C(zipcode) [T.98108]	1.128e+05	1.46e+04	7.738	0.000	8.42e+04
1.41e+05					
C(zipcode) [T.98109]	5.041e+05	1.77e+04	28.512	0.000	4.69e+05
5.39e+05					
C(zipcode) [T.98112]	6.446e+05	1.31e+04	49.298	0.000	6.19e+05

6.7e+05						
C(zipcode) [T.98115]	3.408e+05	1.08e+04	31.486	0.000	3.2e+05	
3.62e+05						
C(zipcode) [T.98116]	2.886e+05	1.23e+04	23.399	0.000	2.64e+05	
3.13e+05						
C(zipcode) [T.98117]	3.208e+05	1.09e+04	29.332	0.000	2.99e+05	
3.42e+05						
C(zipcode) [T.98118]	1.58e+05	1.12e+04	14.158	0.000	1.36e+05	
1.8e+05						
C(zipcode) [T.98119]	4.826e+05	1.47e+04	32.881	0.000	4.54e+05	
5.11e+05						
C(zipcode) [T.98122]	3.484e+05	1.28e+04	27.286	0.000	3.23e+05	
3.73e+05						
C(zipcode) [T.98125]	2.071e+05	1.17e+04	17.742	0.000	1.84e+05	
2.3e+05						
C(zipcode) [T.98126]	1.806e+05	1.21e+04	14.896	0.000	1.57e+05	
2.04e+05						
C(zipcode) [T.98133]	1.682e+05	1.12e+04	15.017	0.000	1.46e+05	
1.9e+05						
C(zipcode) [T.98136]	2.509e+05	1.31e+04	19.107	0.000	2.25e+05	
2.77e+05						
C(zipcode) [T.98144]	2.722e+05	1.22e+04	22.324	0.000	2.48e+05	
2.96e+05						
C(zipcode) [T.98146]	1.026e+05	1.28e+04	8.001	0.000	7.74e+04	
1.28e+05						
C(zipcode) [T.98148]	7.078e+04	2.3e+04	3.075	0.002	2.57e+04	
1.16e+05						
C(zipcode) [T.98155]	1.57e+05	1.14e+04	13.724	0.000	1.35e+05	
1.79e+05						
C(zipcode) [T.98166]	8.926e+04	1.33e+04	6.725	0.000	6.32e+04	
1.15e+05						
C(zipcode) [T.98168]	4.547e+04	1.31e+04	3.472	0.001	1.98e+04	
7.11e+04						
C(zipcode) [T.98177]	2.144e+05	1.33e+04	16.139	0.000	1.88e+05	
2.4e+05						
C(zipcode) [T.98178]	2.97e+04	1.32e+04	2.256	0.024	3900.507	
5.55e+04						
C(zipcode) [T.98188]	2.539e+04	1.63e+04	1.563	0.118	-6459.828	
5.72e+04						
C(zipcode) [T.98198]	9608.5824	1.29e+04	0.745	0.456	-1.57e+04	
3.49e+04						
C(zipcode) [T.98199]	3.836e+05	1.25e+04	30.749	0.000	3.59e+05	
4.08e+05						
bathrooms	5294.7133	2329.863	2.273	0.023	728.009	
9861.418						
sqft_living	154.9575	2.323	66.694	0.000	150.403	
159.512						

```

view                8.339e+04   1581.641    52.723      0.000    8.03e+04
8.65e+04
=====
Omnibus:                19120.795    Durbin-Watson:                1.998
Prob(Omnibus):           0.000    Jarque-Bera (JB):            2895060.214
Skew:                    3.677    Prob(JB):                     0.00
Kurtosis:                59.242    Cond. No.                     1.11e+06
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.11e+06. This might indicate that there are strong multicollinearity or other numerical problems.

"""

So we just got down to five variables and we still got a pretty nice R-squared value of 0.808

## 0.17 trying different models

```

[41]: #preparing our dataframe to perform different models
df4 = df2.copy()
y = df4.price.values # define the target variable (dependent variable) as y
df4.drop('price', axis=1, inplace=True)
zipcode_dummies = pd.get_dummies(df4.zipcode, drop_first=True, prefix = '
↳'zipcode') #introduce dummies for zipcode
grade_dummies = pd.get_dummies(df4.grade, drop_first=True, prefix = 'grade')↳
↳#introduce dummies for grade
dummies = pd.concat([zipcode_dummies, grade_dummies], axis = 1)

```

```

[42]: #drop the features we do not need
df4.drop('sqft_lot', axis=1, inplace=True)
df4.drop('floors', axis=1, inplace=True)
df4.drop('lat', axis=1, inplace=True)
df4.drop('long', axis=1, inplace=True)
df4.drop('zipcode',axis=1, inplace=True)
df4.drop('sqft_living15', axis=1, inplace=True)
df4.drop('sqft_lot15', axis=1, inplace=True)
df4.drop('spring', axis=1, inplace=True)
df4.drop('summer', axis=1, inplace=True)
df4.drop('winter', axis=1, inplace=True)
df4.drop('yr_built', axis=1, inplace=True)
df4.drop('sqft_basement', axis=1, inplace=True)
df4.drop('sqft_above', axis=1, inplace=True)
df4.drop('grade', axis=1, inplace=True)
df4.drop('bedrooms', axis=1, inplace=True)
df4.drop('date', axis=1, inplace=True)

```

```
[43]: #join with our dummies  
df4 = pd.concat([df4, dummies], axis = 1)
```

## 0.18 train/test split

```
[44]: result = []  
  
#split the data  
X_train, X_test, y_train, y_test = train_test_split(df4, y, test_size=0.2)  
  
# fit a model  
lm = linear_model.LinearRegression()  
model = lm.fit(X_train, y_train)  
predictions = lm.predict(X_test)  
result.append(model.score(X_test, y_test))  
print('Train/test split score: ' + str(np.mean(result)))
```

Train/test split score: 0.8279509225077162

## 0.19 K-fold cross validation

```
[45]: # Perform 6-fold cross validation  
scores = cross_val_score(model, df4, y, cv=6)  
print('Cross-validated score:', np.mean(scores))
```

Cross-validated score: 0.824289910931233

## 0.20 Summary

After cleaning and exploration of our data we figured out that the most important feature for house prices in King County. Important features. - location - grade - bathrooms - sqft\_living - view

Considering this features our linear model got a R-squared value of 0.808

## 0.21 Buisness recommendations

House sellers: - try to improve the grade of your house - modelling the fair price of the house to get the maximum

House buyers: - be sure about the location the house is based - calculate the fair price of the house to not pay a premium

## 0.22 Future work

- get more data
- improve data cleaning
  - find a better solution for the cleaning of the waterfront
- get more into detail of the data to find more correlations